

---

# BANK MARKETING ANALYSIS

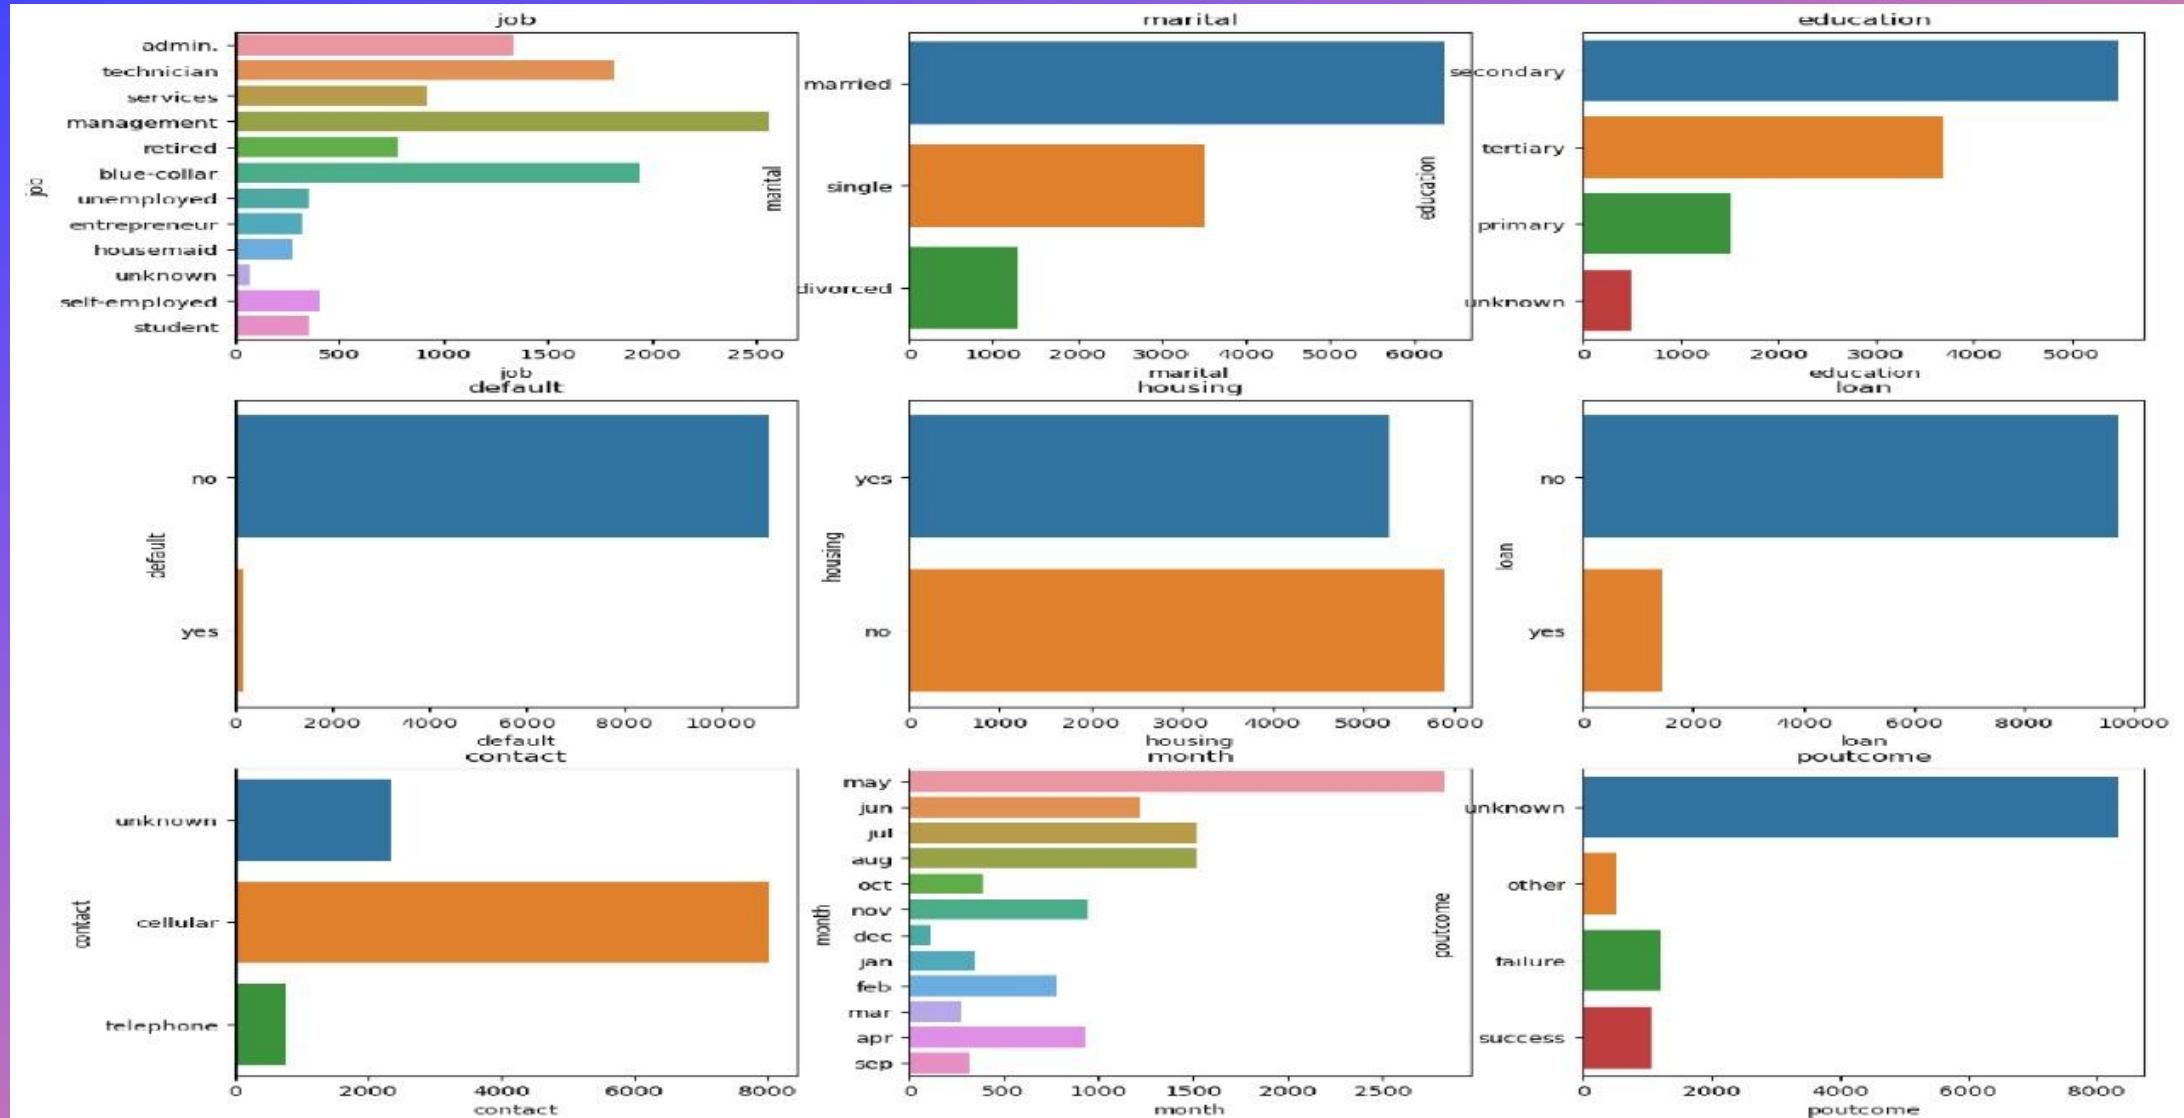
PRESENTED BY:  
AQUILA,KRISHNA

# Overview...



- The Dataset comprises 41,188 rows and 21 columns.
- The dataset contains various features representing client attributes, communication details, and economic indicators.
- Numerical features include age, duration of contact, number of contacts performed during the campaign, and economic factors such as employment variation rate, consumer price index, and consumer confidence index.
- Categorical features encompass client attributes like job type, marital status, education level, and communication channels such as contact type, month, and day of the week of the last contact.
- We will be showcasing implementations in both R and Python for various classification algorithms and feature importance calculation.

# CATEGORICAL VARIABLE VIEW



# DESCRIPTIVE ANALYSIS

## FEATURE DESCRIPTION:

Encompasses client attributes, contact details, economic indicators, and campaign information.

Provides insights into demographic characteristics, communication methods, economic environment, and campaign effectiveness

## IMBALANCED ANALYSIS

Examines the class distribution where one class significantly outnumbers others.

Poses challenges for predictive modeling due to biased algorithms and overlooked minority classes.

## TARGET DEFINITION :

Focuses on the "y" variable indicating term deposit subscription.

Main objective is to accurately predict subscription based on client attributes and campaign interactions.



# CLASSIFICATION

---

Decision Tree, Naive Bayes, Random Forest, Bagging, and Gradient Boosting are classification algorithms used in data analysis.

---

Decision Tree splits data based on attributes, while Naive Bayes is probabilistic.

---

Random Forest combines multiple decision trees, Bagging averages them, and Gradient Boosting sequentially corrects errors.

---

We will be showing Code snippets in R and Python that demonstrate model training and confusion matrix for each algorithm

---

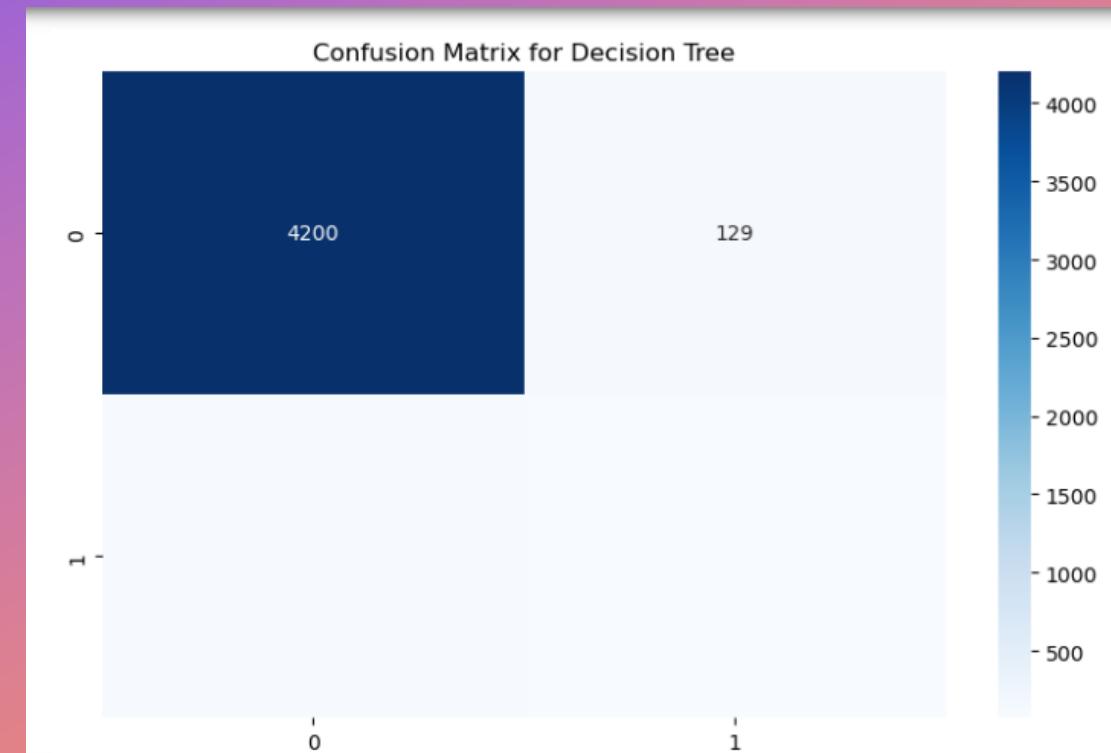
Along with feature selection.

# DECISION TREE

- +
  - o • IMPLEMENTATION IN PYTHON

```
model = DecisionTreeClassifier()  
model.fit(X_train, y_train)
```

```
Statistics for Decision Tree classifier:  
Confusion Matrix:  
[[4192 137]  
 [ 103  68]]  
Accuracy: 0.9466666666666667  
Precision: 0.33170731707317075  
Recall: 0.39766081871345027  
F1 Score: 0.3617021276595745
```



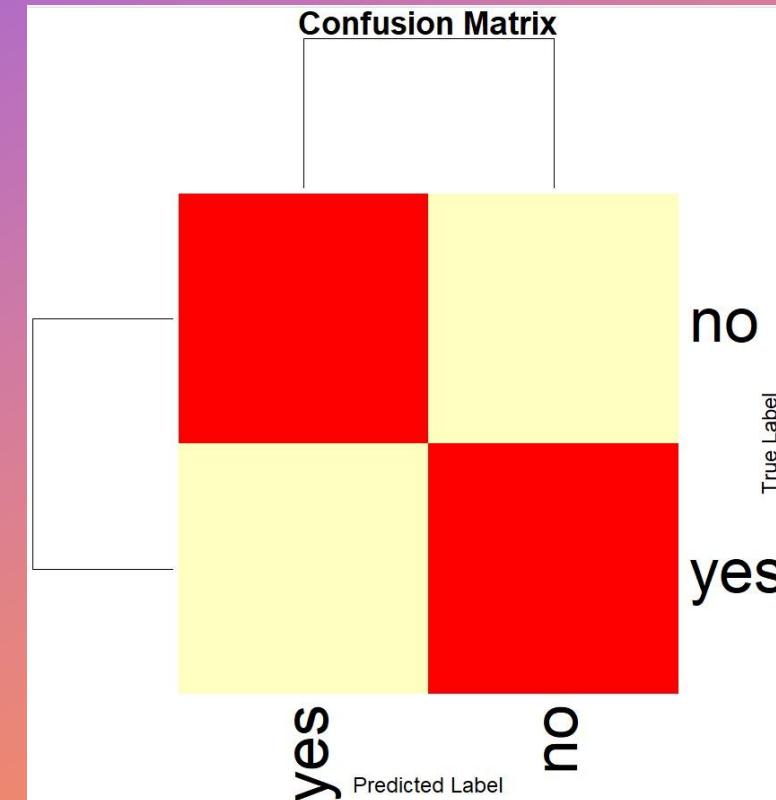
The decision tree is constructed iteratively by selecting the best attribute to split the data at each node.



## IMPLEMENTATION IN R

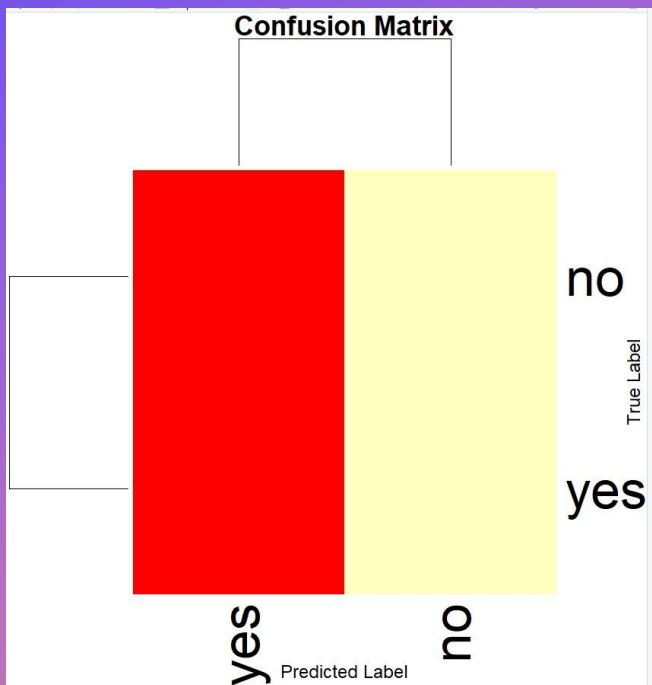
```
# Training Decision Tree model  
model <- rpart(subscription ~ ., data = train_data, method = "class")
```

```
+ }  
> # Iterate over models and train/evaluate  
> for (model_name in names(models_cat)) {  
+   train_and_evaluate_cat_model(model_name, models_cat[[model_name]], train_data_  
at, test_data_cat)  
+ }  
Confusion Matrix for Decision Tree :  
Confusion Matrix and Statistics  
  
Reference  
Prediction no yes  
no 4239 90  
yes 79 90  
  
Accuracy : 0.9624  
95% CI : (0.9565, 0.9678)  
No Information Rate : 0.96  
P-Value [Acc > NIR] : 0.2134  
  
Kappa : 0.4962  
  
McNemar's Test P-Value : 0.4418  
  
Sensitivity : 0.9817  
Specificity : 0.5000  
Pos Pred Value : 0.9792  
Neg Pred Value : 0.5325  
Prevalence : 0.9600  
Detection Rate : 0.9424  
Detection Prevalence : 0.9624  
Balanced Accuracy : 0.7409  
  
'Positive' Class : no
```



# NAÏVE BAYES

## IMPLEMENTATION IN R



```
model <- naive_bayes(subscription ~ ., data = train_data)
```

Confusion Matrix for Naive Bayes :  
Confusion Matrix and Statistics

Reference  
Prediction    no    yes  
              no    4318    180  
              yes    0    0

Accuracy : 0.96  
95% CI : (0.9538, 0.9655)  
No Information Rate : 0.96  
P-Value [Acc > NIR] : 0.5198

Kappa : 0

McNemar's Test P-Value : <2e-16

Sensitivity : 1.00  
Specificity : 0.00  
Pos Pred Value : 0.96  
Neg Pred Value : NaN  
Prevalence : 0.96  
Detection Rate : 0.96  
Detection Prevalence : 1.00  
Balanced Accuracy : 0.50

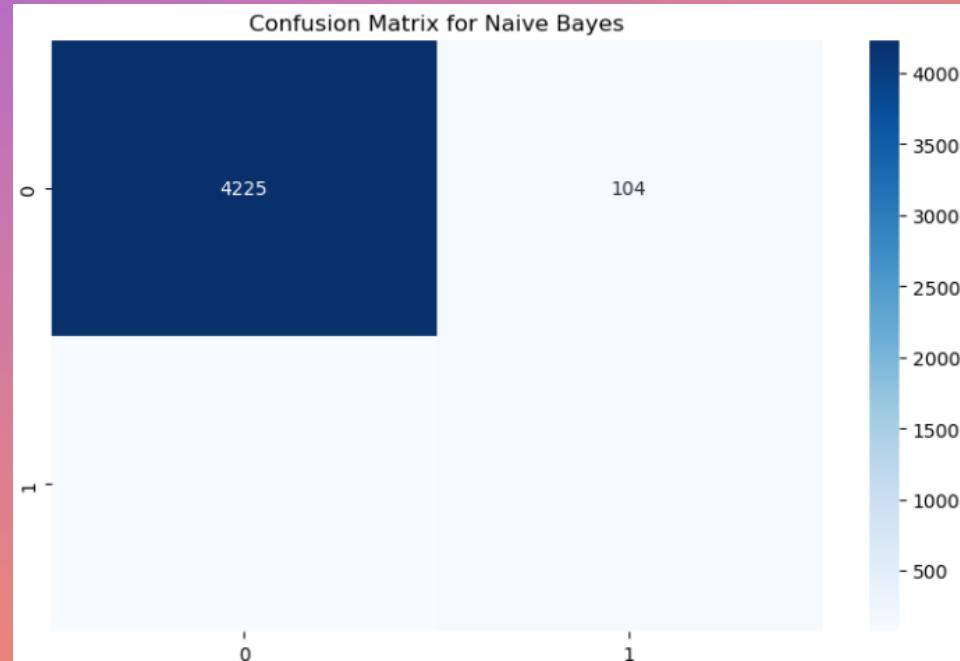
'Positive' Class : no

# IMPLEMENTATION IN PYTHON

IT ASSUMES THAT THE FEATURES ARE CONDITIONALLY INDEPENDENT GIVEN THE CLASS, HENCE THE "NAIVE" ASSUMPTION

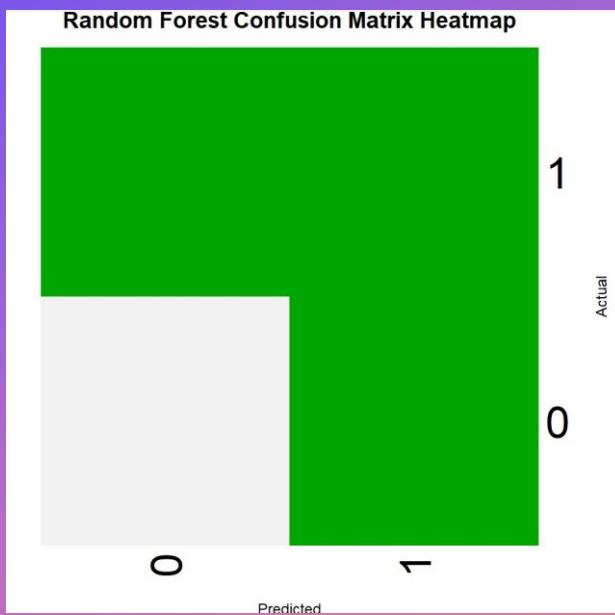
```
model = GaussianNB()  
model.fit(X_train, y_train)
```

Statistics for Naive Bayes classifier:  
Confusion Matrix:  
[[4225 104]  
 [ 68 103]]  
Accuracy: 0.9617777777777777  
Precision: 0.4975845410628019  
Recall: 0.6023391812865497  
F1 Score: 0.544973544973545



# RANDOM FOREST

## IMPLEMENTATION IN R.

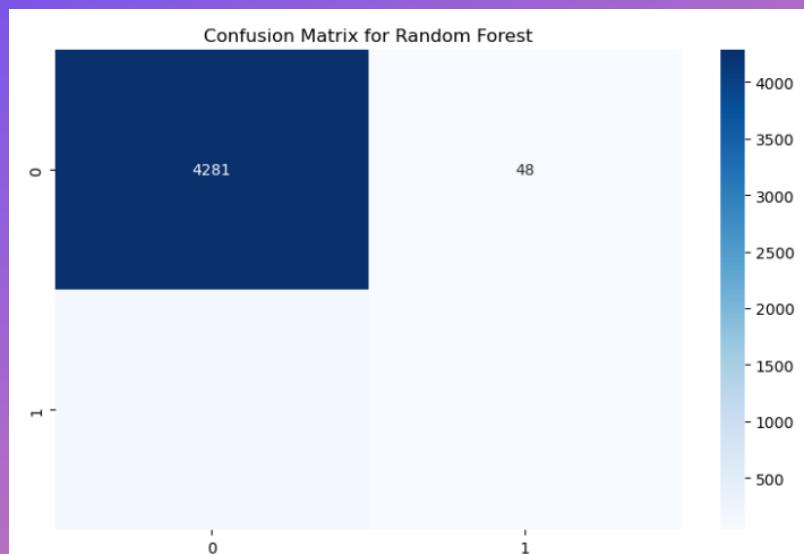


```
rf_model <- randomForest(subscription~, data=train, ntree=50,  
                           ntry=6, importance=TRUE, replace=FALSE)
```

```
> confusionMatrix (original, rf_pred)  
Confusion Matrix and Statistics  
  
Reference  
Prediction    0     1  
          0 4785    24  
          1   136    16  
  
Accuracy : 0.9677  
95% CI  : (0.9624, 0.9725)  
No Information Rate : 0.9919  
P-Value [Acc > NIR] : 1  
  
Kappa : 0.1559  
  
McNemar's Test P-Value : <2e-16  
  
Sensitivity : 0.9724  
Specificity  : 0.4000  
Pos Pred Value : 0.9950  
Neg Pred Value : 0.1053  
Prevalence   : 0.9919  
Detection Rate : 0.9645  
Detection Prevalence : 0.9694  
Balanced Accuracy : 0.6862  
  
'Positive' Class : 0
```

# IMPLEMENTATION IN PYTHON

**IT CONSTRUCTS MULTIPLE DECISION TREES DURING TRAINING AND OUTPUTS THE MODE OF THE CLASSES (CLASSIFICATION) OR MEAN PREDICTION (REGRESSION) OF THE INDIVIDUAL TREES**



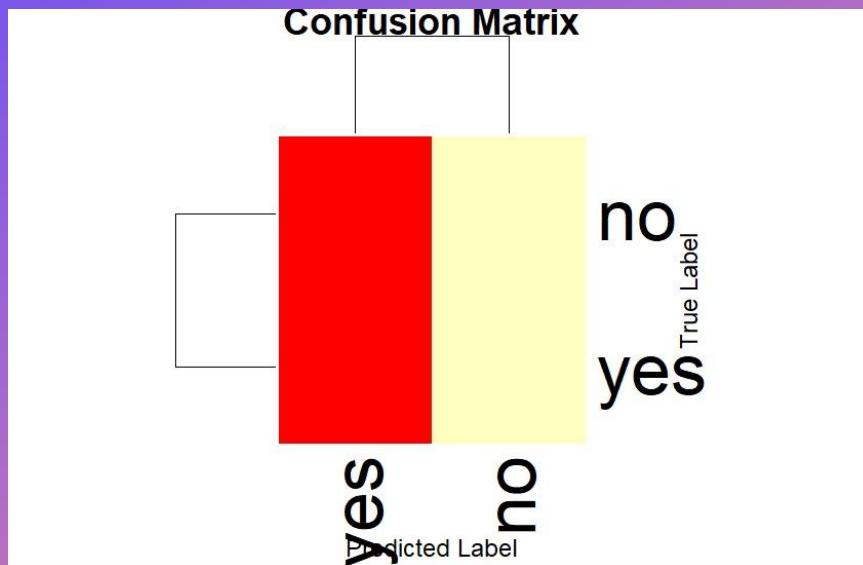
```
model = RandomForestClassifier(n_estimators=100)  
model.fit(X_train, y_train)
```

```
Statistics for Random Forest classifier:  
Confusion Matrix:  
[[4283    46]  
 [ 133    38]]  
Accuracy: 0.9602222222222222  
Precision: 0.4523809523809524  
Recall: 0.2222222222222222  
F1 Score: 0.2980392156862745
```

# GRADIENT BOOSTING

## IMPLEMENTATION IN R

```
model <- abm(subscription ~ ., data = train_data, n.trees = 100, interaction.depth = 4)
```



```
9      0.1930      nan  0.1000  0.0014
10     0.1882      nan  0.1000  0.0021
20     0.1666      nan  0.1000  0.0007
40     0.1576      nan  0.1000 -0.0000
50     0.1551      nan  0.1000 -0.0001

Confusion Matrix for Gradient Boosting :
Confusion Matrix and Statistics

Reference
Prediction  no  yes
      no  4272  122
      yes   46   58

Accuracy : 0.9627
95% CI : (0.9567, 0.968)
No Information Rate : 0.96
P-Value [Acc > NIR] : 0.1915

Kappa : 0.3906

Mcnemar's Test P-Value : 7.192e-09

Sensitivity : 0.9893
Specificity : 0.3222
Pos Pred Value : 0.9722
Neg Pred Value : 0.5577
Prevalence : 0.9600
Detection Rate : 0.9498
Detection Prevalence : 0.9769
Balanced Accuracy : 0.6558

'Positive' Class : no
```

# IMPLEMENTATION IN PYTHON

+

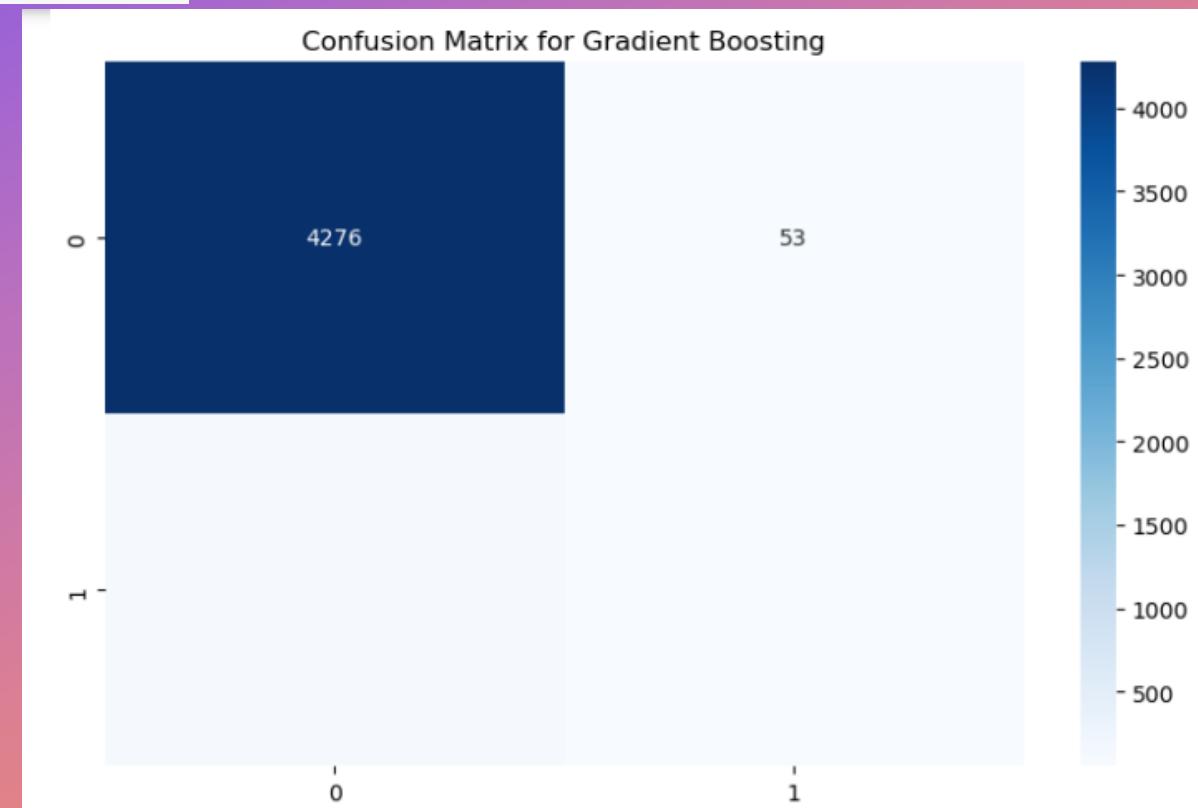
.

Gradient Boosting is effective for a wide range of predictive modeling tasks and is known for its high predictive accuracy

```
model = GradientBoostingClassifier(n_estimators=100)  
model.fit(X_train, y_train)
```

o

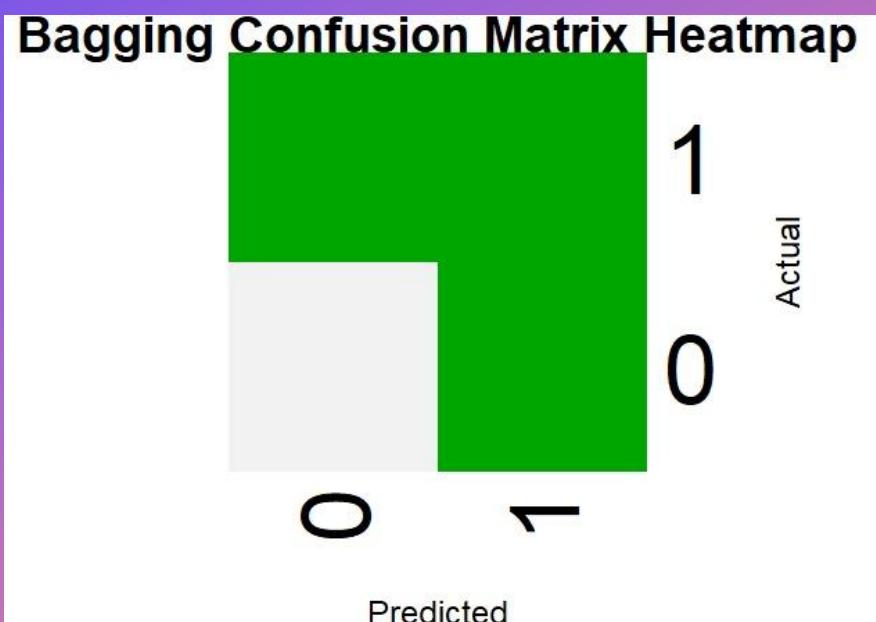
```
Statistics for Gradient Boosting classifier:  
Confusion Matrix:  
[[4276  53]  
 [ 13  58]]  
Accuracy: 0.9631111111111111  
Precision: 0.5225225225225225  
Recall: 0.3391812865497076  
F1 Score: 0.41134751773049644
```



# BAGGING

+  
•  
o

## IMPLEMENTATION IN R



```
bag_model <- bagging(formula = subscription~, data = train,  
                      nbagg=50, coob=TRUE)
```

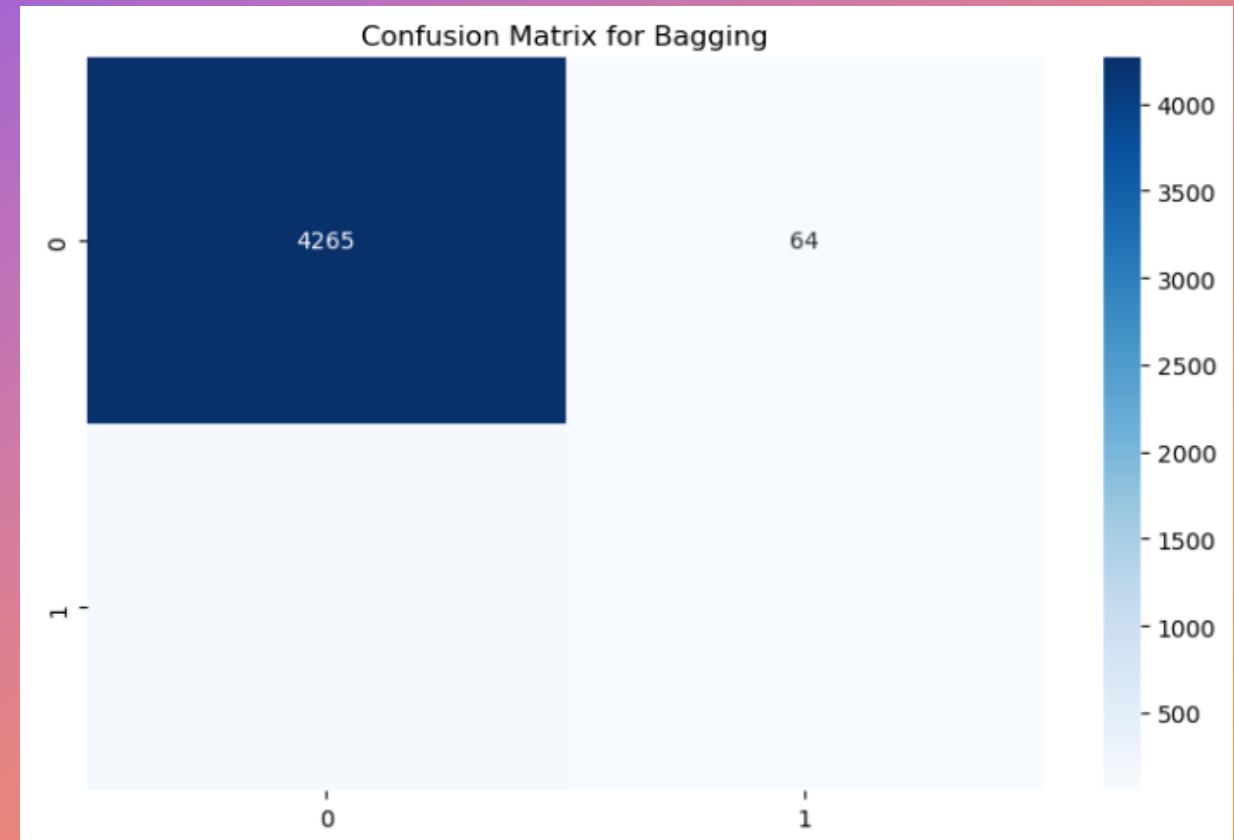
```
> confusionMatrix (original, bag_pred)  
Confusion Matrix and Statistics  
  
Reference  
Prediction 0 1  
0 4745 64  
1 103 49  
  
Accuracy : 0.9663  
95% CI : (0.9609, 0.9712)  
No Information Rate : 0.9772  
P-Value [Acc > NIR] : 0.999999  
  
Kappa : 0.3529  
  
McNemar's Test P-Value : 0.003277  
  
Sensitivity : 0.9788  
Specificity : 0.4336  
Pos Pred Value : 0.9867  
Neg Pred Value : 0.3224  
Prevalence : 0.9772  
Detection Rate : 0.9565  
Detection Prevalence : 0.9694  
Balanced Accuracy : 0.7062  
  
'Positive' Class : 0
```

# IMPLEMENTATION IN PYTHON

It Builds multiple models independently and then combines their predictions through averaging (for regression) or voting.

```
model <- bagging(subscription ~ ., data = train_data, nbagg = 50)
```

```
Statistics for Bagging classifier:  
Confusion Matrix:  
[[4263  66]  
 [ 126  45]]  
Accuracy: 0.9573333333333334  
Precision: 0.40540540540540543  
Recall: 0.2631578947368421  
F1 Score: 0.3191489361702128
```



# IMPLEMENTATION IN R

# SUPPORT VECTOR MACHINE

```
svm1 <- svm(subscription~, data = train, kernel = "radial")
```

Confusion Matrix and Statistics

Reference		Prediction	
		0	1
0	11626	855	
1	524	724	

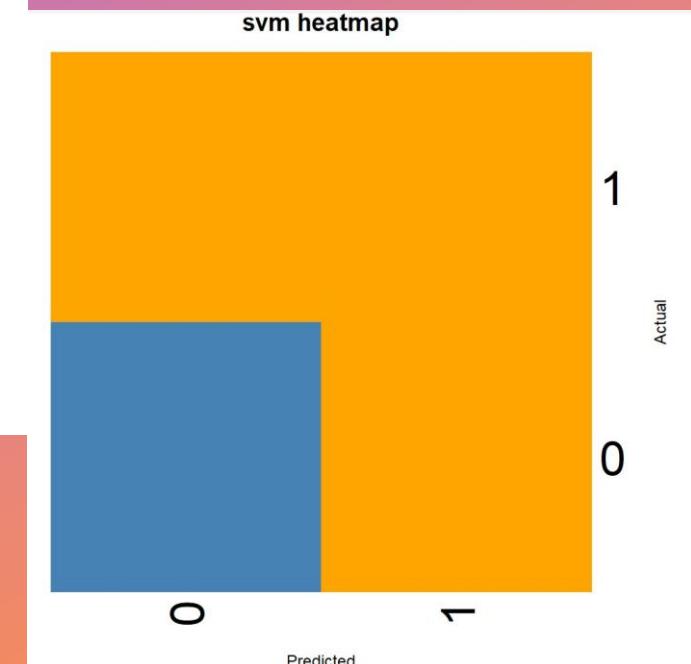
Accuracy : 0.8996  
95% CI : (0.8944, 0.9045)  
No Information Rate : 0.885  
P-Value [Acc > NIR] : 2.711e-08

Kappa : 0.4571

McNemar's Test P-Value : < 2.2e-16

Sensitivity : 0.9569  
Specificity : 0.4585  
Pos Pred Value : 0.9315  
Neg Pred Value : 0.5801  
Prevalence : 0.8850  
Detection Rate : 0.8468  
Detection Prevalence : 0.9091  
Balanced Accuracy : 0.7077

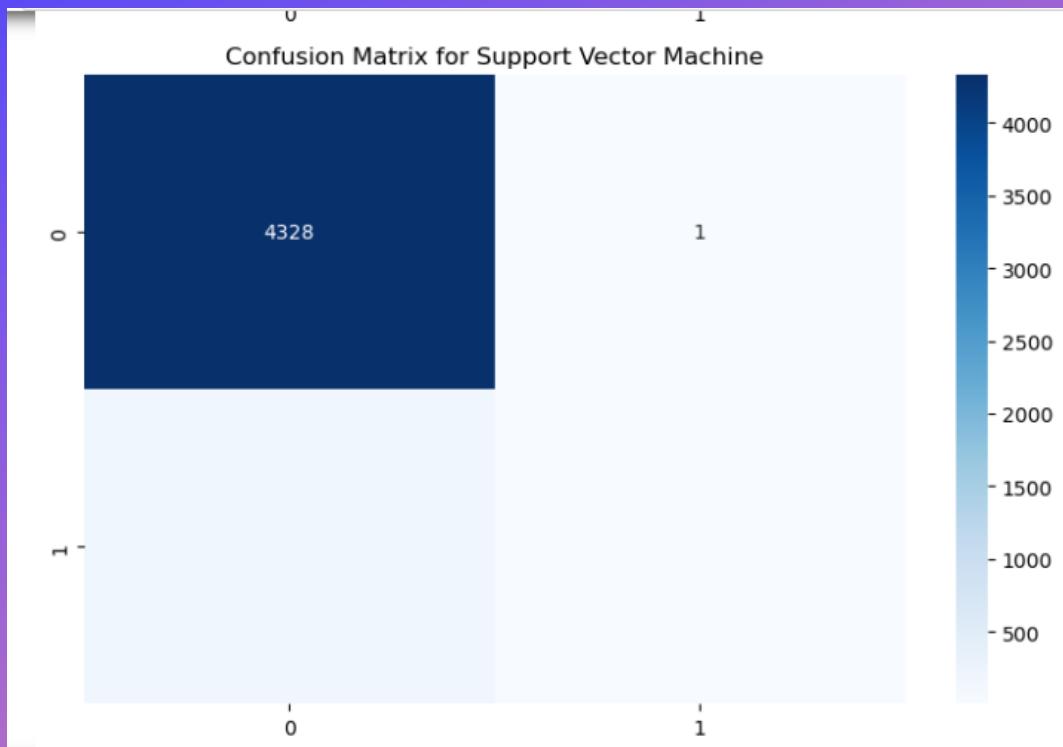
'Positive' Class : 0



# IMPLEMENTATION IN PYTHON

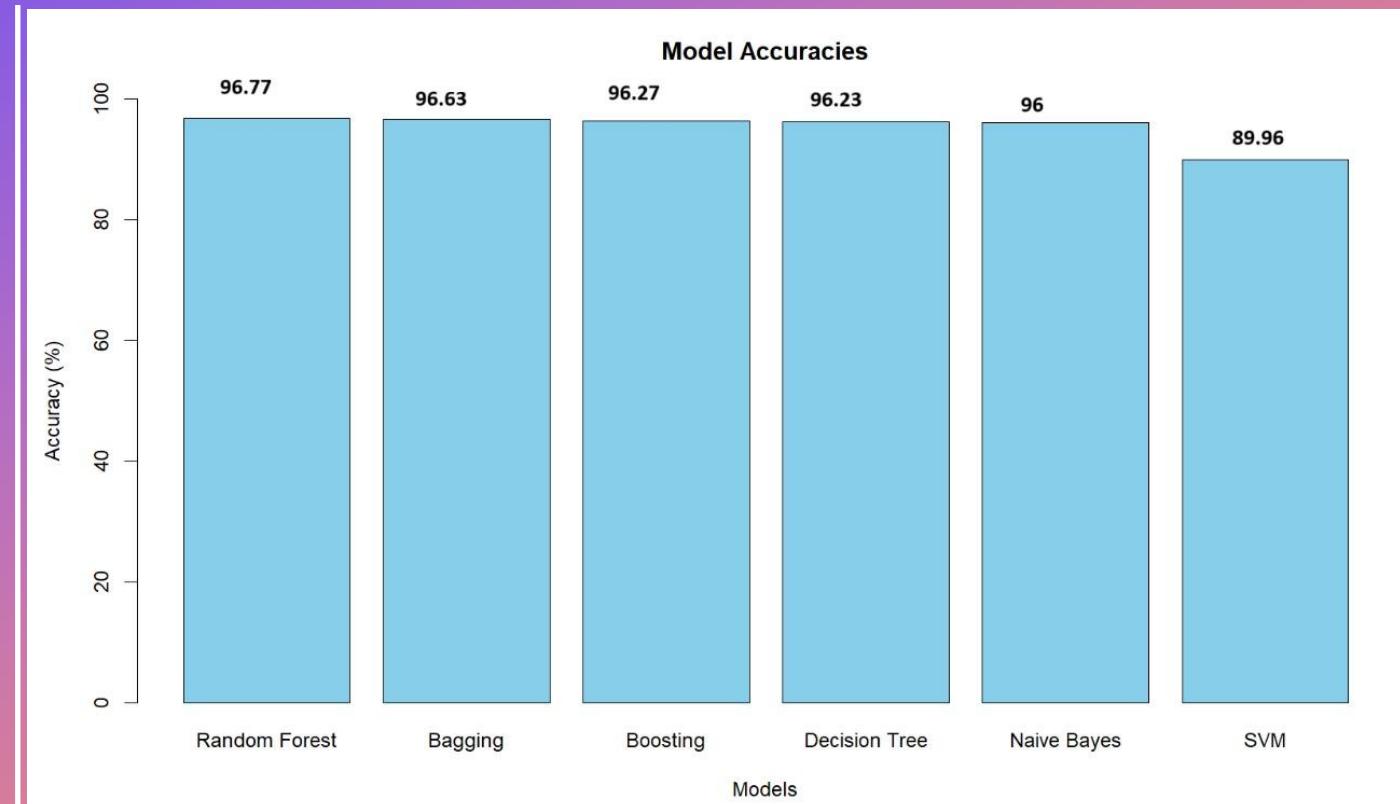
SVM aims to maximize the margin between classes, making it robust to outliers and effective in high-dimensional spaces.

```
"Support Vector Machine": SVC()  
svm_classifier.fit(X_train, y_train)
```



```
Statistics for Support Vector Machine classifier:  
Confusion Matrix:  
[[4328  1]  
 [ 168   3]]  
Accuracy: 0.9624444444444444  
Precision: 0.75  
Recall: 0.017543859649122806  
F1 Score: 0.03428571428571429
```

# COMPARISON OF ALL CATEGORICAL MODELS



# NUMERICAL VALUES CLASSIFICATION



## METHODS USED IN R

KNN - The 'k' nearest data points, determined by a distance metric such as Euclidean distance, are considered for classification, We Implemented this classification with various k values and found the best accuracy for k=22

NAÏVE BAYES – We have tried this specific method for comparison of the accuracy with other numerical categories.

## METHODS USED IN PYTHON

KNN

NAÏVE BAYES

DECISION TREE

This iterates over various k values (ODD NUMBERS – 3,5,7) for KNN and creates multiple instances of Decision Tree and Naive Bayes classifiers. By training and evaluating these models on the same dataset, it enables direct comparison of their performance metrics, aiding in selecting the most suitable algorithm for the classification task.

# KNN

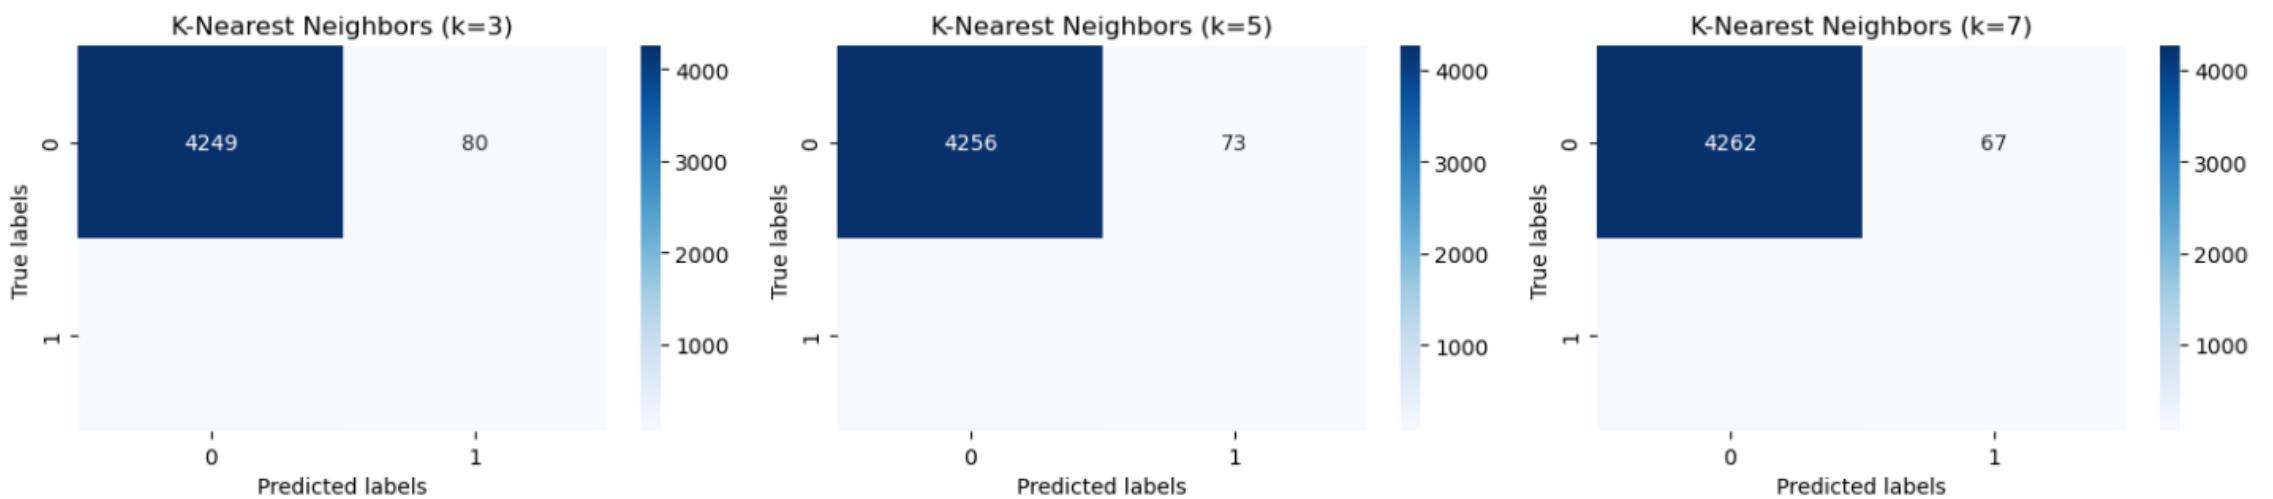
## IN PYTHON

```
classifiers = {  
    "K-Nearest Neighbors": [KNeighborsClassifier(n_neighbors=k) for k in k_values],
```

Statistics for K-Nearest Neighbors (Model 1):  
Confusion Matrix:  
[[4249 80]  
 [ 112 59]]  
Accuracy: 0.9573333333333334  
Precision: 0.4244604316546763  
Recall: 0.34502923976608185  
F1 Score: 0.38064516129032255

Statistics for K-Nearest Neighbors (Model 2):  
Confusion Matrix:  
[[4256 73]  
 [ 112 59]]  
Accuracy: 0.9588888888888889  
Precision: 0.44696969696969696  
Recall: 0.34502923976608185  
F1 Score: 0.38943894389438943

Statistics for K-Nearest Neighbors (Model 3):  
Confusion Matrix:  
[[4262 67]  
 [ 116 55]]  
Accuracy: 0.9593333333333334  
Precision: 0.45081967213114754  
Recall: 0.3216374269005848  
F1 Score: 0.37542662116040953



# NAÏVE BAYES IN PYTHON

Statistics for Naive Bayes (Model 1):

Confusion Matrix:

```
[[4225 104]
 [ 68 103]]
```

Accuracy: 0.9617777777777777  
Precision: 0.4975845410628019  
Recall: 0.6023391812865497  
F1 Score: 0.544973544973545

Statistics for Naive Bayes (Model 2):

Confusion Matrix:

```
[[4225 104]
 [ 68 103]]
```

Accuracy: 0.9617777777777777  
Precision: 0.4975845410628019  
Recall: 0.6023391812865497  
F1 Score: 0.544973544973545

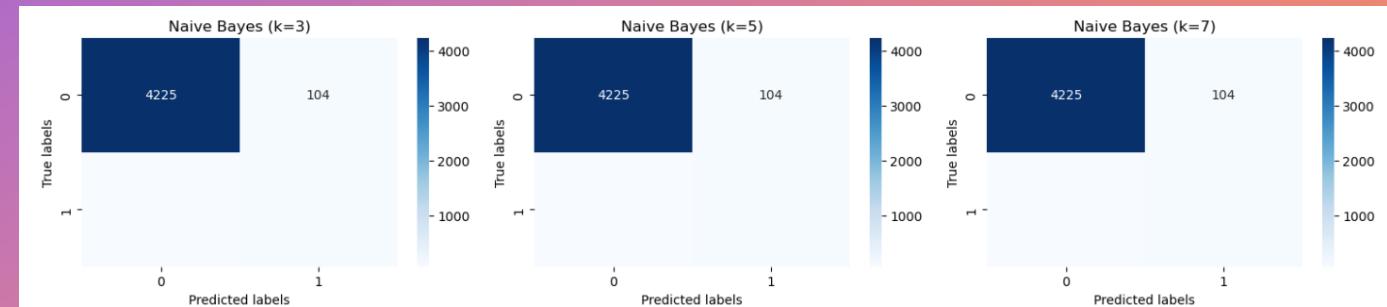
Statistics for Naive Bayes (Model 3):

Confusion Matrix:

```
[[4225 104]
 [ 68 103]]
```

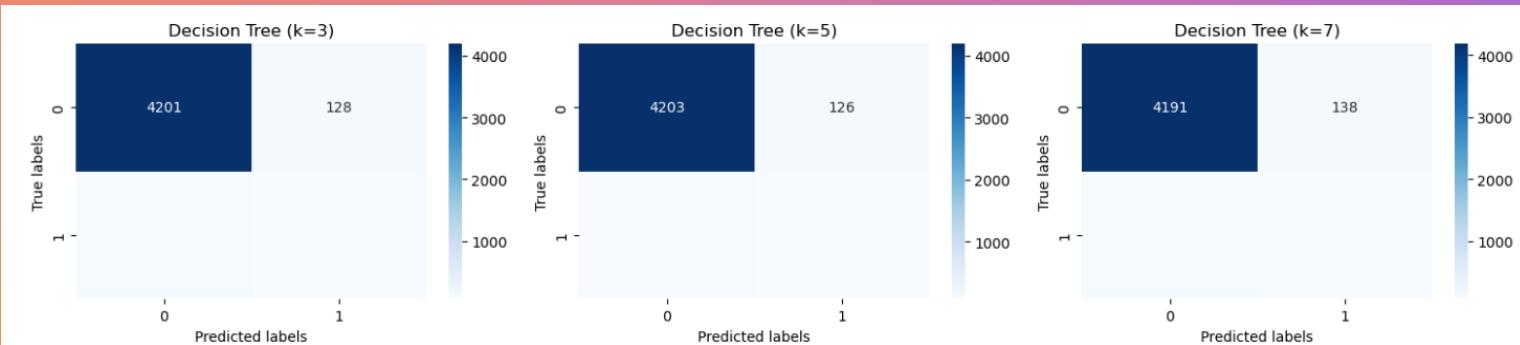
Accuracy: 0.9617777777777777  
Precision: 0.4975845410628019  
Recall: 0.6023391812865497  
F1 Score: 0.544973544973545

```
classifiers = {}
"Naive Bayes": [GaussianNB() for _ in range(len
```



```
classifiers = {  
    "Decision Tree": [DecisionTreeClassifier() for _ in range(len
```

# DECISION TREE IN PYTHON



Statistics for Decision Tree (Model 1):  
Confusion Matrix:  
[[4195 134]  
[ 109 62]]  
Accuracy: 0.946  
Precision: 0.3163265306122449  
Recall: 0.36257309941520466  
F1 Score: 0.33787465940054495

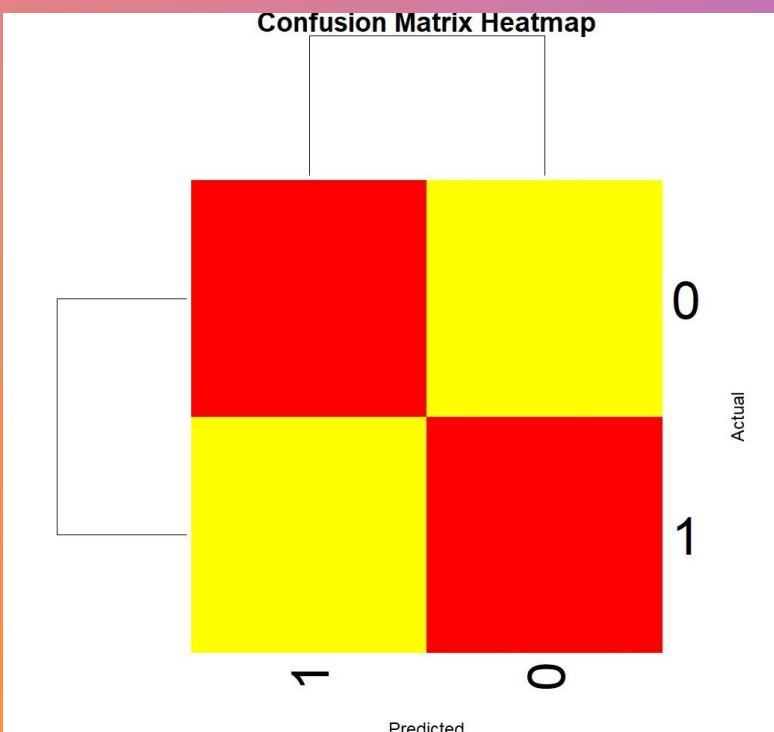
Statistics for Decision Tree (Model 2):  
Confusion Matrix:  
[[4197 132]  
[ 107 64]]  
Accuracy: 0.9468888888888889  
Precision: 0.32653061224489793  
Recall: 0.3742690058479532  
F1 Score: 0.34877384196185285

Statistics for Decision Tree (Model 3):  
Confusion Matrix:  
[[4207 122]  
[ 103 68]]  
Accuracy: 0.95  
Precision: 0.35789473684210527  
Recall: 0.39766081871345027  
F1 Score: 0.3767313019390582

# KNN

## IMPLEMENTATION IN R

```
# Fitting KNN Model to training dataset
classifier_knn <- knn(train = train_scale, test = test_scale,
                      cl = train$subscription, k = 5)
classifier_knn <- knn(train = train_scale, test = test_scale,
                      cl = train$subscription, k = 10)
classifier_knn <- knn(train = train_scale, test = test_scale,
                      cl = train$subscription, k = 22)
```



```
> confusionMatrix(classifier_knn, actual_labels)
Confusion Matrix and Statistics

Reference
Prediction      0      1
      0 11798    844
      1    374    713

Accuracy : 0.9113
95% CI  : (0.9064, 0.916)
No Information Rate : 0.8866
P-Value [Acc > NIR] : < 2.2e-16

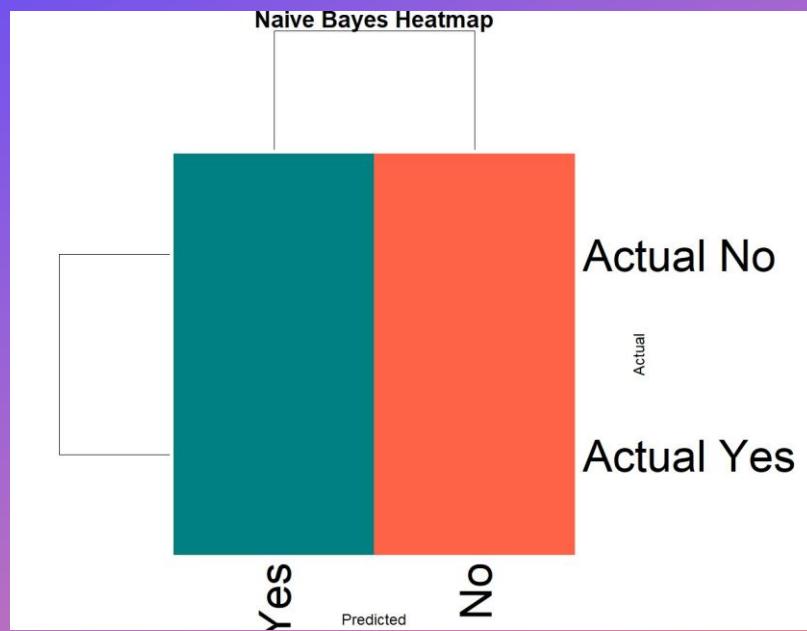
Kappa : 0.492

McNemar's Test P-Value : < 2.2e-16

Sensitivity : 0.9693
Specificity  : 0.4579
Pos Pred Value : 0.9332
Neg Pred Value : 0.6559
Prevalence   : 0.8866
Detection Rate : 0.8593
Detection Prevalence : 0.9208
Balanced Accuracy : 0.7136

'Positive' Class : 0
```

# NAÏVE BAYES IMPLEMENTATION IN R



```
# Fitting Naive Bayes Model to training dataset  
classifier_nb <- naiveBayes(x = train_scale, y = train$subscription)
```

## Confusion Matrix and Statistics

		Reference	
		0	1
Prediction	0	11067	634
	1	1105	923

Accuracy : 0.8733  
95% CI : (0.8677, 0.8789)  
No Information Rate : 0.8866  
P-Value [Acc > NIR] : 1

Kappa : 0.4435

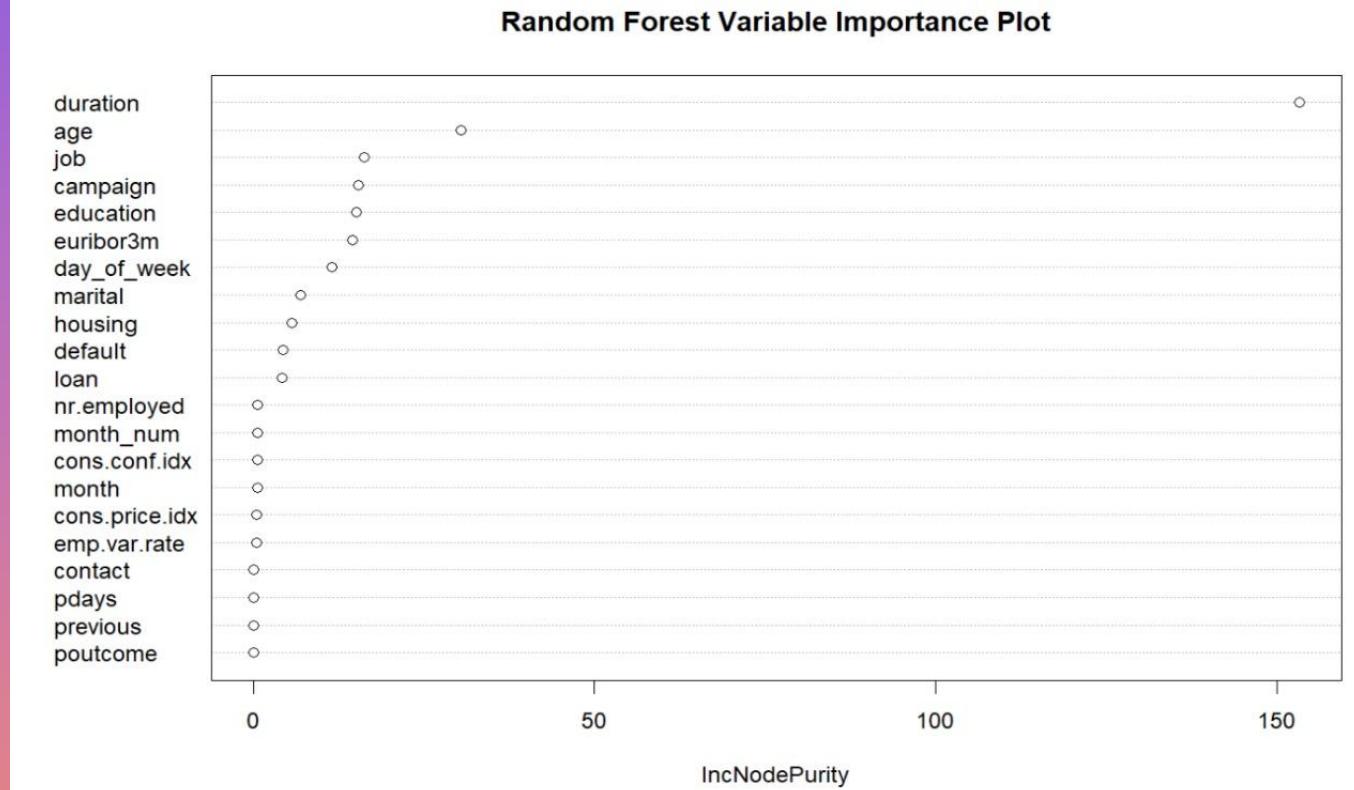
McNemar's Test P-Value : <2e-16

Sensitivity : 0.9092  
Specificity : 0.5928  
Pos Pred Value : 0.9458  
Neg Pred Value : 0.4551  
Prevalence : 0.8866  
Detection Rate : 0.8061  
Detection Prevalence : 0.8523  
Balanced Accuracy : 0.7510

'Positive' Class : 0

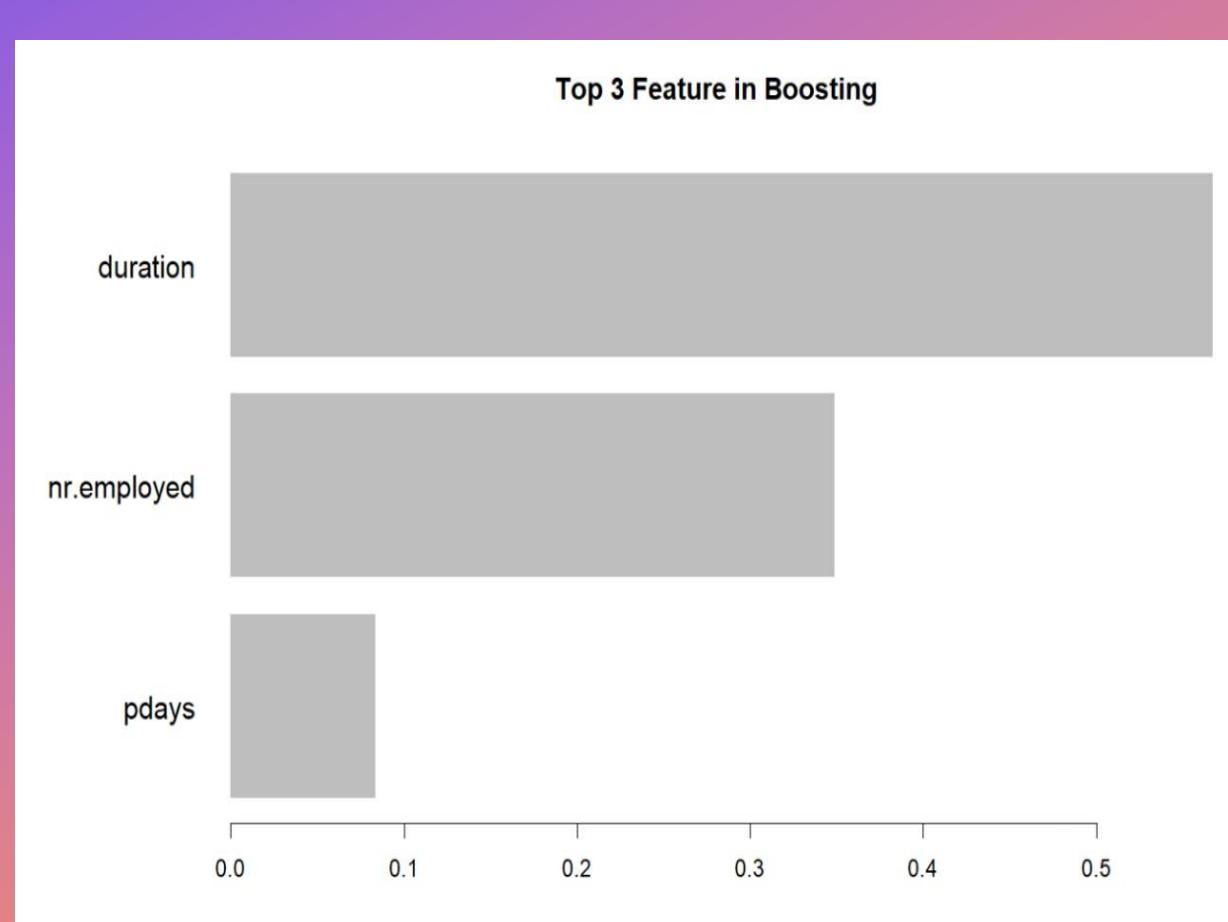
# FEATURE IMPORTANCE CALCULATION USING RANDOM FOREST METHOD

```
# Fitting Random Forest Model  
rf_model <- randomForest(subscription ~ ., data = bank_data)  
# Plotting variable importance  
varImpPlot(rf_model, main = "Random Forest Variable Importance Plot")
```



# FEATURE IMPORTANCE CALCULATION USING BOOSTING

```
# Plotting feature importance
importance_matrix <- xgb.importance(feature_names = colnames(train_data[, -1]),
                                      model = boost_model)
xgb.plot.importance(importance_matrix[1:3], main = "Top 3 Feature in Boosting")
```



# CLUSTERING

```
# Identify categorical and numerical variables  
categorical_vars <- data[, sapply(data, is.factor)]  
numerical_vars <- data[, sapply(data, is.numeric)]  
  
# Create datasets for categorical and numerical variables  
categorical_data <- data[, sapply(data, is.factor)]  
numerical_data <- data[, sapply(data, is.numeric)]
```

- K MEANS - KMeans clustering is a popular unsupervised machine learning algorithm used for clustering data points into a predefined number of clusters. The algorithm aims to partition the data into clusters such that the similarity within each cluster is maximized, while the similarity between clusters is minimized.

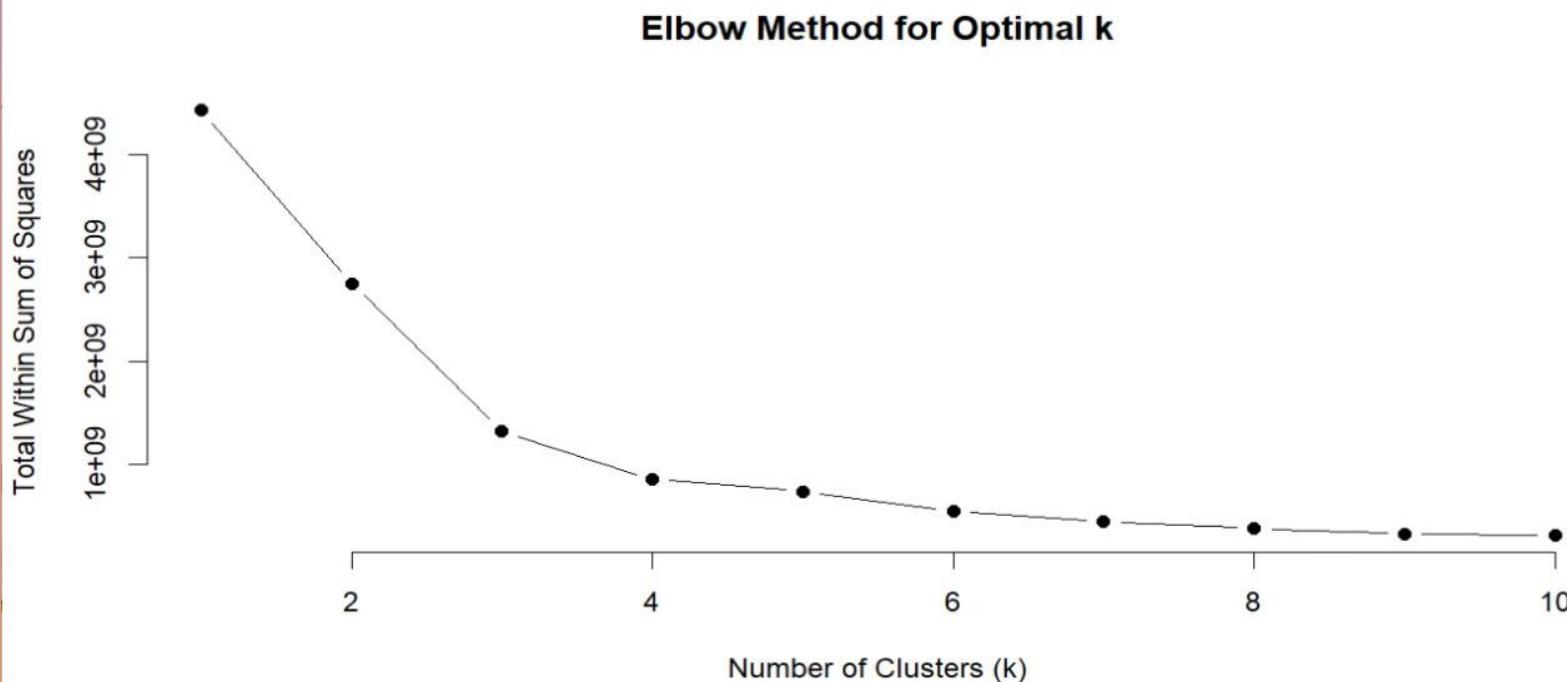
```

# Function to calculate total within-cluster sum of squares
wss <- function(k) {
  kmeans_model <- kmeans(numerical_data, centers = k)
  return(sum(kmeans_model$withinss))
}

# Compute within-cluster sum of squares for different values of k
k_values <- 1:10
wss_values <- sapply(k_values, wss)

# Plot the elbow curve
plot(k_values, wss_values, type = "b", pch = 19, frame = FALSE,
      xlab = "Number of Clusters (k)", ylab = "Total Within Sum of Squares",
      main = "Elbow Method for Optimal k")

```



```
#K means for clusters = 3 from elbow curve  
kmeans_model <- kmeans(numerical_data, 3)  
kmeans_model
```

```
R 4.3.2 · ~/R Studio Works/ALY 6040/Final Project/ ↵
> kmeans_model <- kmeans(numerical_data, 3)
> kmeans_model
K-means clustering with 3 clusters of sizes 34859, 1509, 4820
```

```

Cluster means:
      age duration campaign days_after_contacted_p previous previous_interest_rate
1 39.94765 180.5454 2.617201          999.0000 0.1181331           0.1662842
2 41.86547 314.5427 1.821074            6.0000 1.6626905          -2.0964877
3 40.00021 802.8973 2.442531          997.7685 0.1031120           0.1534855
      cons.price.idx cons.conf.idx euribor3m nr.employed      y
1         93.58295     -40.55035 3.7252459    5172.459 0.05020224
2         93.34247     -38.33227 0.9859768    5029.251 0.63817097
3         93.59599     -40.83676 3.6945102    5170.954 0.39979253

```

Clustering vector:

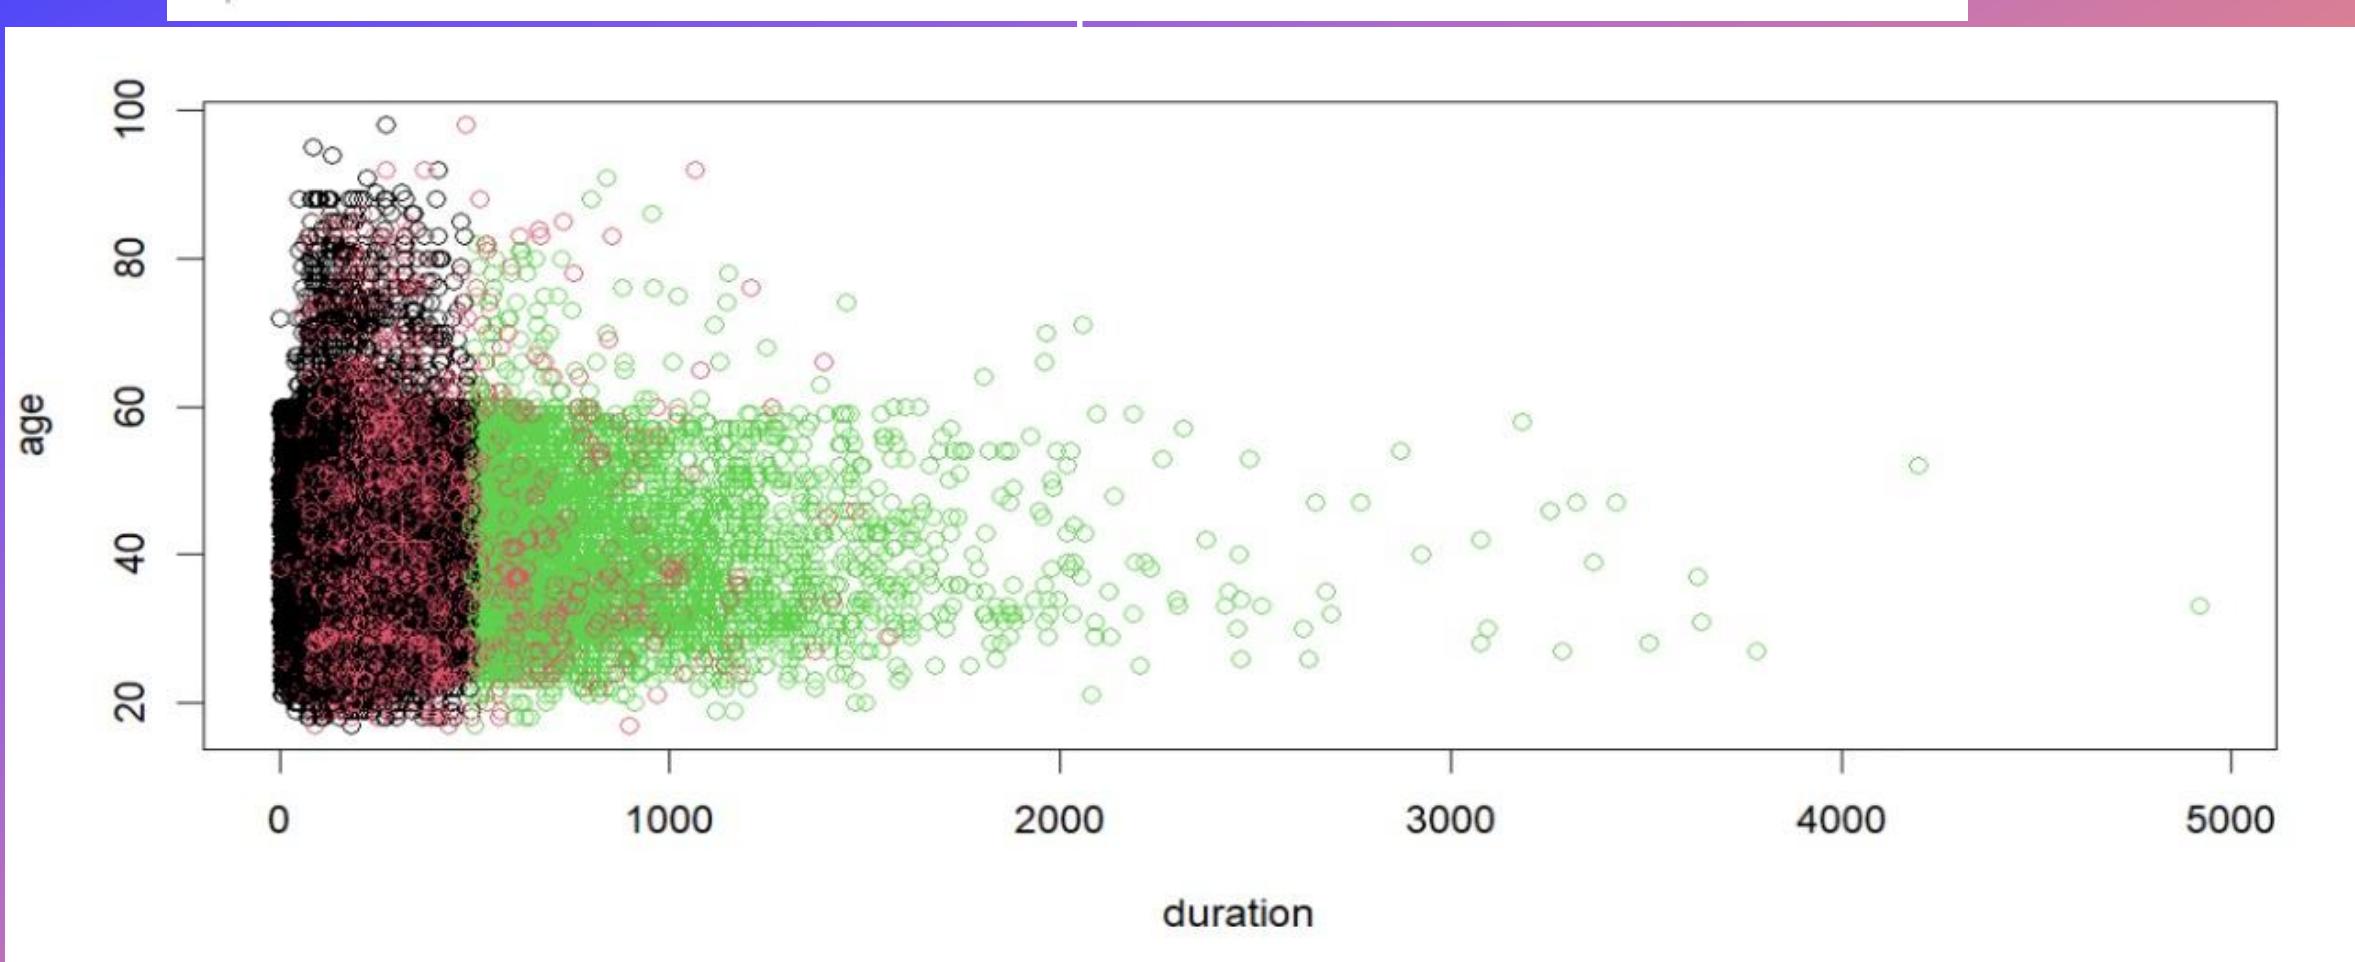
```
> table(numerical_data$y, kmeans_model$cluster)
```

	1	2	3
0	33109	546	2893
1	1750	963	1927

```
> plot(numerical_data[c("duration", "age")], col=kmeans_model$cluster)
```

```
> points(kmeans_model$centers[, c("duration", "age")], col=1:3, pch=8, cex=2)
```

```
>
```



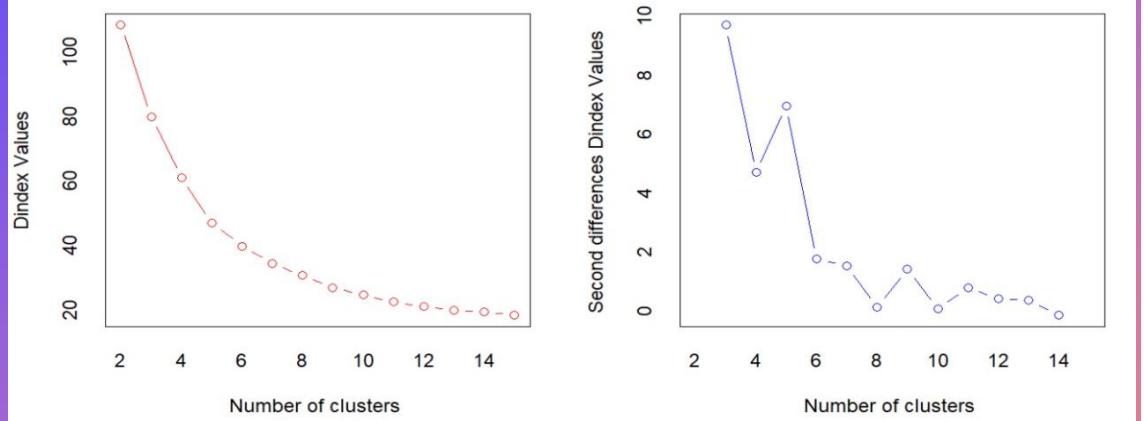
# NB CLUST METHOD

```
library(cluster)
library(NbClust)

# Perform NbClust analysis
numerical_data_subset <- numerical_data[1:1000,1:3]

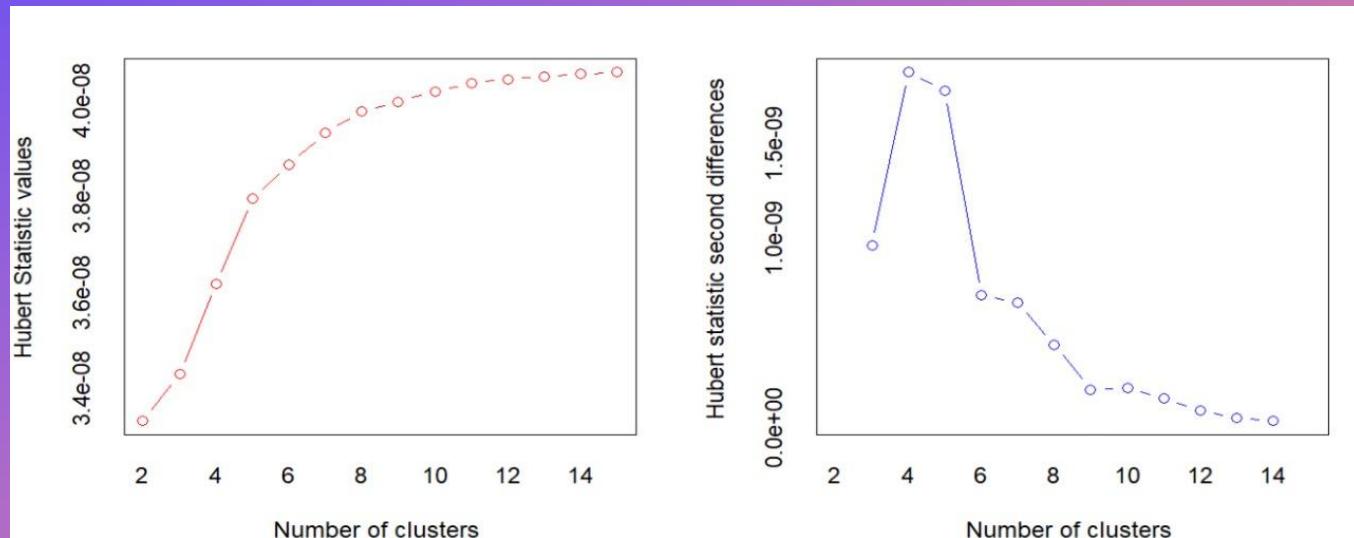
library(caret)

# Perform cluster analysis
nbclust_result <- NbClust(data = numerical_data_subset, min.r)
Bar<-table(nbclust_result$Best.n[1,])
Bar
```



```
> Bar<-table(nbclust_result$Best.n[1,])
> Bar
```

```
0 2 3 4 6 9 11 15
2 8 4 1 1 1 7 2
```



```
> nbclust_result <- NbClust(data = numerical_data_subset, min.nc = 2, max.nc = 15, method = "kmeans")
*** : The Hubert index is a graphical method of determining the number of clusters.
In the plot of Hubert index, we seek a significant knee that corresponds to a
significant increase of the value of the measure i.e the significant peak in Hubert
index second differences plot.
```

```
*** : The D index is a graphical method of determining the number of clusters.
In the plot of D index, we seek a significant knee (the significant peak in Dindex
second differences plot) that corresponds to a significant increase of the value of
the measure.
```

```
*****
* Among all indices:
```

- \* 8 proposed 2 as the best number of clusters
- \* 4 proposed 3 as the best number of clusters
- \* 1 proposed 4 as the best number of clusters
- \* 1 proposed 6 as the best number of clusters
- \* 1 proposed 9 as the best number of clusters
- \* 7 proposed 11 as the best number of clusters
- \* 2 proposed 15 as the best number of clusters

\*\*\*\*\* Conclusion \*\*\*\*\*

\* According to the majority rule, the best number of clusters is 2

# IMPLEMENTATION IN PYTHON

```
: data['y'] = data['y'].apply(lambda x: 1 if x == 'yes' else 0)

: sns.pairplot(data, hue="y", palette="viridis", plot_kws={"alpha": 0.4})
plt.show()
```



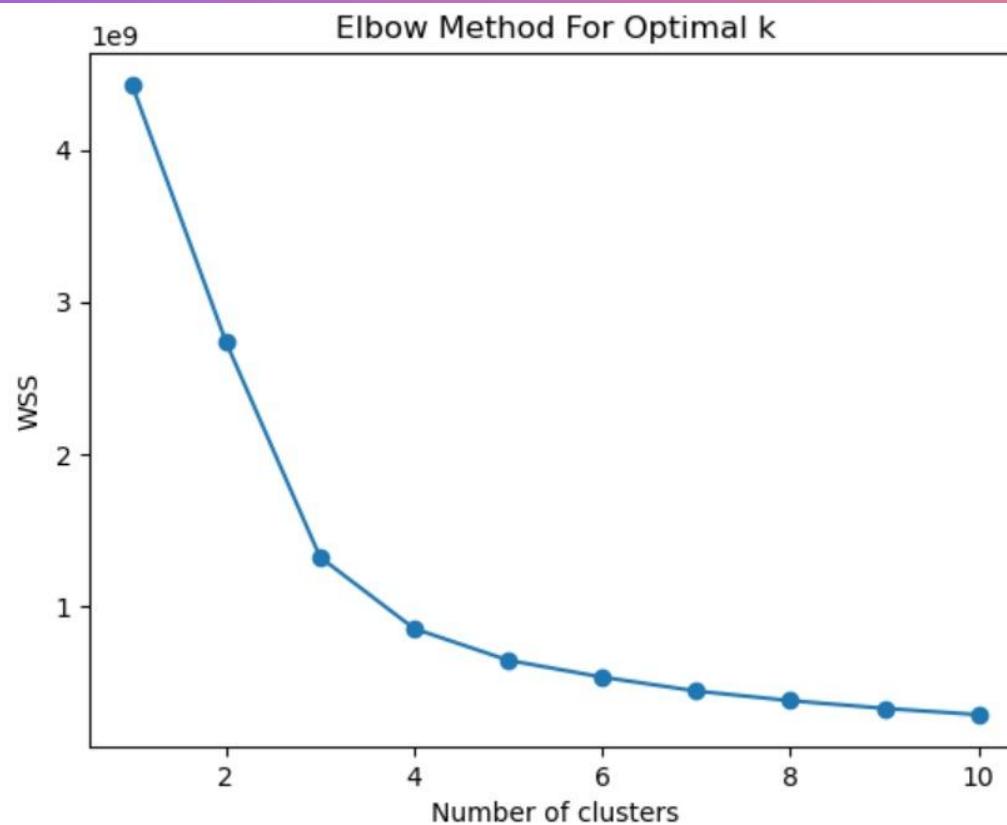
# WITHIN CLUSTER SUM OF SQUARE

```
k_opt = 3
kmeans_opt = KMeans(n_clusters=k_opt, random_state=42)
kmeans_opt.fit(numerical_data)
print("Cluster means:")
print(kmeans_opt.cluster_centers_)

C:\Users\trill\anaconda3\Lib\site-packages\sklearn\cluster\_kmeans.py:870: FutureWarning: The default value of `n_init` will change from 4 to 10 in 0.23. Set the value of `n_init` explicitly to suppress the warning
  warnings.warn(
Cluster means:
[[ 3.99981351e+01  8.02509532e+02  2.44135930e+00  9.97769996e+02
   1.03398259e-01  1.50497306e-01  9.35950253e+01 -4.08389142e+01
   3.69140489e+00  5.17084279e+03  3.99709905e-01]
 [ 3.99479241e+01  1.80491981e+02  2.61739305e+00  9.99000000e+02
   1.18096003e-01  1.66700141e-01  9.35830800e+01 -4.05499986e+01
   3.72568120e+00  5.17247435e+03  5.01535019e-02]
 [ 4.18654738e+01  3.14542744e+02  1.82107356e+00  6.00000000e+00
   1.66269052e+00 -2.09648774e+00  9.33424679e+01 -3.83322730e+01
   9.85976806e-01  5.02925070e+03  6.38170974e-01]]
```

```
: wss = []
for k in range(1, 11):
    kmeans = KMeans(n_clusters=k, n_init=10, random_state=42) # Explicitly set n_init
    kmeans.fit(numerical_data)
    wss.append(kmeans.inertia_)

: plt.plot(range(1, 11), wss, marker='o')
plt.title('Elbow Method For Optimal k')
plt.xlabel('Number of clusters')
plt.ylabel('WSS')
plt.show()
```



# SILHOUETTE METHOD

```
: k_opt = 3
kmeans_opt = KMeans(n_clusters=k_opt, random_state=42)
kmeans_opt.fit(numerical_data)
print("Cluster means:")
print(kmeans_opt.cluster_centers_)

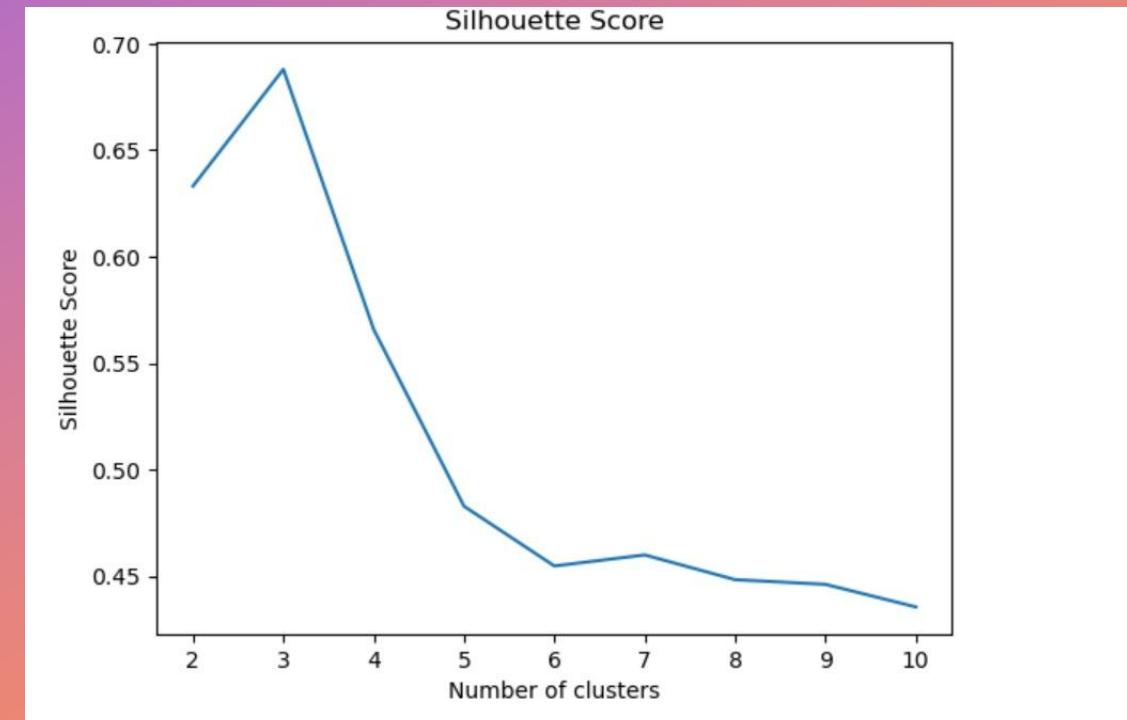
C:\Users\trill\anaconda3\Lib\site-packages\sklearn\cluster\_kmeans.py:870: FutureWarning: The default value of `n_init` will change from 10 to 14. Set the value of `n_init` explicitly to suppress the warning
  warnings.warn(
Cluster means:
[[ 3.99981351e+01  8.02509532e+02  2.44135930e+00  9.97769996e+02
  1.03398259e-01  1.58497306e-01  9.35950253e+01 -4.08389142e+01
  3.69140489e+00  5.17084279e+03  3.99709905e-01]
 [ 3.99479241e+01  1.80491981e+02  2.61739305e+00  9.99000000e+02
  1.18096003e-01  1.66700141e-01  9.35830800e+01 -4.05499986e+01
  3.72568120e+00  5.17247435e+03  5.01535019e-02]
 [ 4.18654738e+01  3.14542744e+02  1.821073556e+00  6.00000000e+00
  1.66269052e+00 -2.09648774e+00  9.33424679e+01 -3.83322730e+00
  9.85976806e-01  5.02925070e+03  6.38170974e-01]]
```

```
: # Silhouette Score
silhouette_scores = []
for n_clusters in range(2, 11):
    kmeans = KMeans(n_clusters=n_clusters, init='k-means++', max_iter=300, n_init=10, random_state=0)
    cluster_labels = kmeans.fit_predict(numerical_data)
    silhouette_avg = silhouette_score(numerical_data, cluster_labels)
    silhouette_scores.append(silhouette_avg)
silhouette = silhouette_scores

[0.6331208146929761, 0.6879613893051137, 0.5658051326698728, 0.48280792718349985, 0.45476621438284814, 0.45994780696935533, 0.4483395658146482, 0.4461410
425019763, 0.435545925189039]

print(silhouette)

plt.plot(range(2, 11), silhouette_scores)
plt.title('Silhouette Score')
plt.xlabel('Number of clusters')
plt.ylabel('Silhouette Score')
plt.show()
```



# IMPLEMENTATION IN R

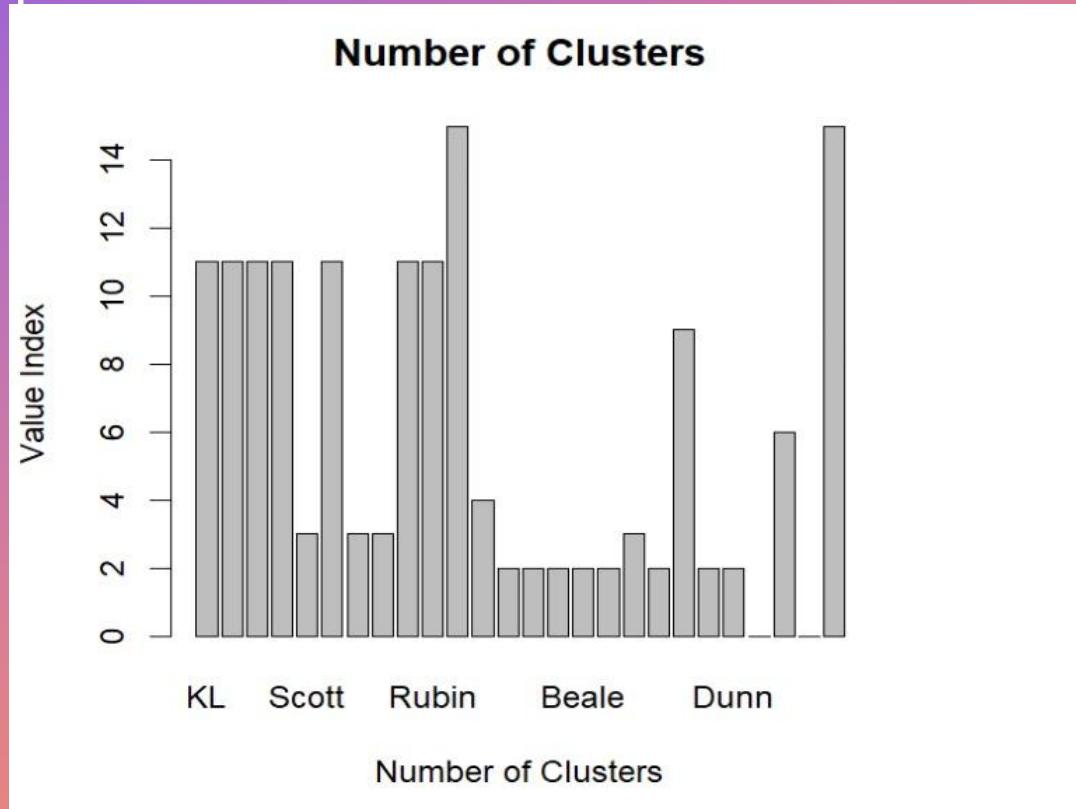
# PARTITIONING AROUND MEDOIDS WITH K-MEDOIDS

```
numerical_data_subset <- numerical_data[1:1000,1:3 ]
nbclust_result <- NbClust(data = numerical_data_subset, min.nc = 2, max.nc = 15, meth

str(nbclust_result$Best.nc)

# Extract the number of clusters from NbClust results
num_clusters <- nbclust_result$Best.nc[1, ]

# Create a bar plot
barplot(num_clusters, main = "Number of Clusters",
        xlab = "Number of Clusters", ylab = "Value Index")
```



# IN PYTHON

```
import numpy as np
import matplotlib.pyplot as plt
from sklearn_extra.cluster import KMedoids
from sklearn.metrics import silhouette_score

silhouette_array = np.array(silhouette)
# Calculate silhouette scores for different numbers of clusters
n_clusters_range = range(2, 11)
silhouette_scores = []
```

```
for n_clusters in n_clusters_range:
    if n_clusters <= len(silhouette): # Check if the number of clusters is valid
        kmedoids = KMedoids(n_clusters=n_clusters, random_state=42)
        cluster_labels = kmedoids.fit_predict(silhouette.reshape(-1, 1))

        if len(np.unique(cluster_labels)) < 2 or len(np.unique(cluster_labels)) >= len(silhouette):
            silhouette_avg = np.nan # Set silhouette score to NaN
        else:
            silhouette_avg = silhouette_score(silhouette.reshape(-1, 1), cluster_labels)

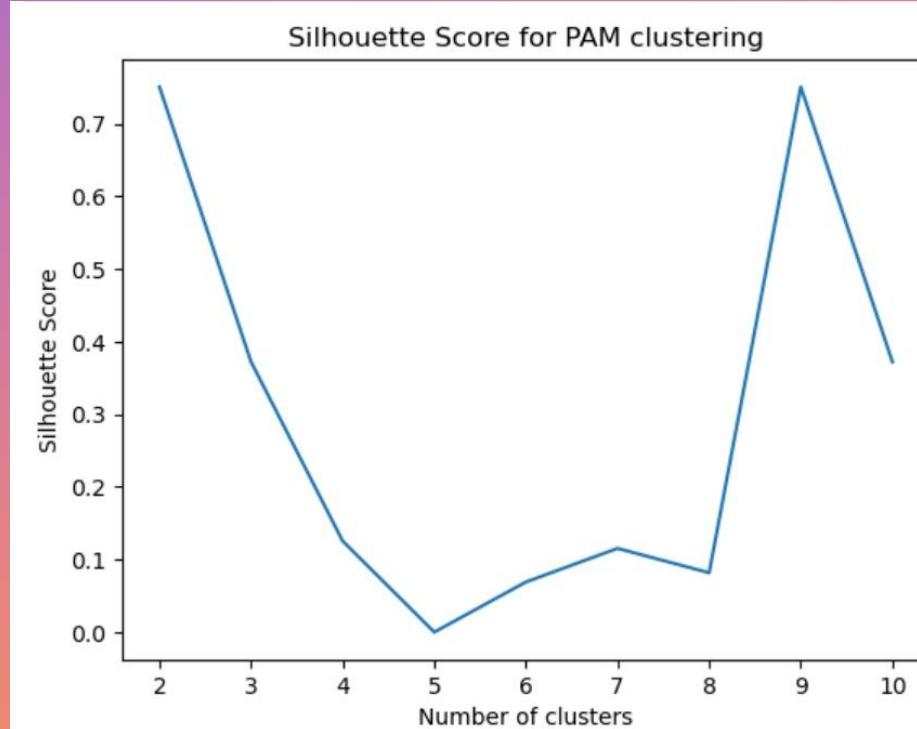
        silhouette_scores.append(silhouette_avg)

    # Print silhouette scores
print(silhouette_scores)

[0.7504423966663645, 0.3720195503389193, 0.12546674711425723, 0.00014474006572864288, 0.06910587800231835, 0.1153747745128902, 0.08182048973199998, 0.750
44239666663645, 0.3720195503389193, 0.12546674711425723, 0.00014474006572864288, 0.06910587800231835, 0.1153747745128902, 0.08182048973199998, nan, nan]
```

```
n_clusters_range = range(2, 11)

# Plot silhouette scores corresponding to the range of clusters
plt.plot(n_clusters_range, silhouette_scores[:len(n_clusters_range)])
plt.title('Silhouette Score for PAM clustering')
plt.xlabel('Number of clusters')
plt.ylabel('Silhouette Score')
plt.show()
```



# DENSITY BASED SPATIAL CLUSTERING OF APPLICATIONS WITH NOISE

+

```
install.packages("dbSCAN")
library(dbSCAN)

# Perform DBSCAN clustering
dbSCAN_result <- dbSCAN(numerical_data_subset, eps = 0.5, minPts = 5)

# Extract cluster assignments and noise points
cluster_assignments <- dbSCAN_result$cluster
noise_points <- numerical_data_subset$dbSCAN_result$cluster == 0, ]

# Analyze the clusters and noises
num_clusters <- length(unique(cluster_assignments)) - 1 # Subtract 1 for noise cluster
num_noises <- sum(cluster_assignments == 0)
```

```
> # Print the number of clusters and noises
> cat("Number of clusters:", num_clusters, "\n")
Number of clusters: 3
> cat("Number of noise points:", num_noises, "\n")
Number of noise points: 0
> |
```

# IN PYTHON

```
from sklearn.cluster import DBSCAN
from sklearn.metrics import silhouette_score
import numpy as np

# Input data: silhouette values
silhouette = np.array([0.63312081, 0.68796139, 0.56580513, 0.48280793, 0.45476621, 0.45994781, 0.44833957, 0.4461

# Reshape silhouette data to fit DBSCAN input format
X = silhouette.reshape(-1, 1)

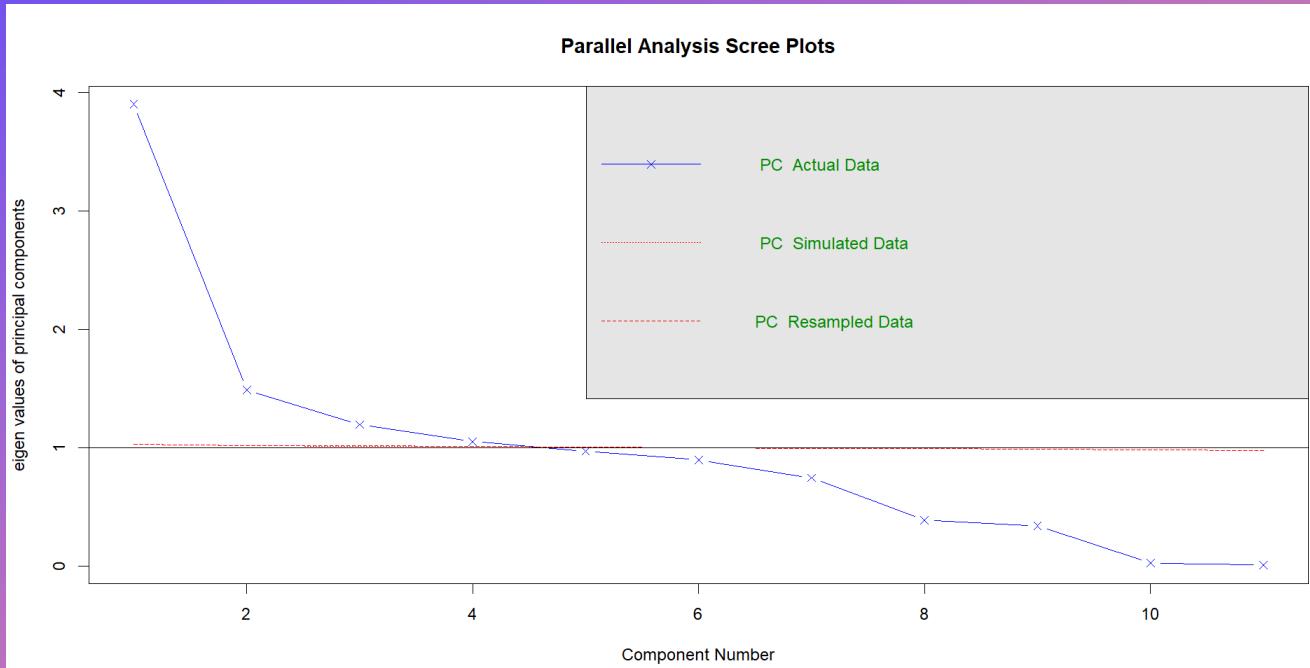
# DBSCAN clustering
dbscan = DBSCAN(eps=0.1, min_samples=2)
cluster_labels = dbscan.fit_predict(X)

# Analyzing clusters and noises
n_clusters_ = len(set(cluster_labels)) - (1 if -1 in cluster_labels else 0) # Number of clusters, ignoring noise
n_noises_ = list(cluster_labels).count(-1) # Number of noise points

print('Estimated number of clusters:', n_clusters_)
print('Estimated number of noise points:', n_noises_)
print('Cluster labels:', cluster_labels)
```

# PCA IMPLEMENTATION IN R

```
# PCA Scree plot  
fa.parallel(bank_data[,-1], fa="pc", n.iter=100, show.legend = TRUE)
```



We are choosing 4 as best component because that is where the elbow point of the curve is and the curve starts to flatten out.

# CREATING DATA FRAME FOR 4 COMPONENTS AND APPLYING RF MODEL TO FIND THE ACCURACY

```
> confusionMatrix(table(rf_pred, test_labels))
Confusion Matrix and Statistics

      test_labels
rf_pred   0     1
  0 8815  574
  1  370  538

      Accuracy : 0.9083
      95% CI  : (0.9026, 0.9138)
No Information Rate : 0.892
P-Value [Acc > NIR] : 2.667e-08

      Kappa : 0.4824

McNemar's Test P-Value : 3.920e-11

      Sensitivity : 0.9597
      Specificity : 0.4838
      Pos Pred Value : 0.9389
      Neg Pred Value : 0.5925
      Prevalence : 0.8920
      Detection Rate : 0.8561
      Detection Prevalence : 0.9118
      Balanced Accuracy : 0.7218

'Positive' Class : 0
```

```
+ # Perform PCA
  pca_model <- prcomp(bank_data[, -1], scale = TRUE)
  pca <- as.data.frame(pca_model$x[, 1:4])
```

```
• # Random Forest
  rf_model <- randomForest(subscription ~ ., data =
    train_data_with_sub, ntree = 50)
```

# CONCLUSION

This project analyzed a bank marketing dataset employing a comprehensive range of machine learning techniques including Decision Trees, SVM, Random Forest, Bagging, Boosting, Naive Bayes for categorical data, KNN, Random Forest, and Naive Bayes for numerical data, alongside clustering.

Insights gleaned facilitated personalized marketing campaigns, leveraging customer segmentation, predictive modeling, and ensemble learning. Recommendations emphasize continuous model refinement and integration into marketing strategies for enhanced customer engagement and conversion. By embracing data-driven methodologies, the project aims to empower the bank with actionable insights to optimize marketing efforts, bolster customer relationships, and drive sustainable growth in the financial services sector.

# REFERENCES

- 
- Moro, S., Cortez, P., & Rita, P. (2014). A Data-Driven Approach to Predict the Success of Bank Telemarketing. *Decision Support Systems*.
  - Zanella, A., Olsina, L., & Salto, C. (2017). Predictive Models for Bank Telemarketing Campaigns. *Expert Systems with Applications*.
  - Silva, A., & Ribeiro, R. (2018). Predictive Modeling in Direct Marketing: An Application Using Data Mining Techniques. *Journal of Business Research*.